

Deep Reinforcement Learning for Delay-Optimized Task Offloading in Vehicular Fog Computing

Mohammad Parsa Toopchinezhad - Mahmood Ahmadi

Presentation Outline

I

Introduction

- The Internet of Vehicles
- Task Offloading
- Reinforcement Learning

II

Problem Formulation

- Vehicle Model
- Task Model
- Communication Model
- MDP Specification

III

Results

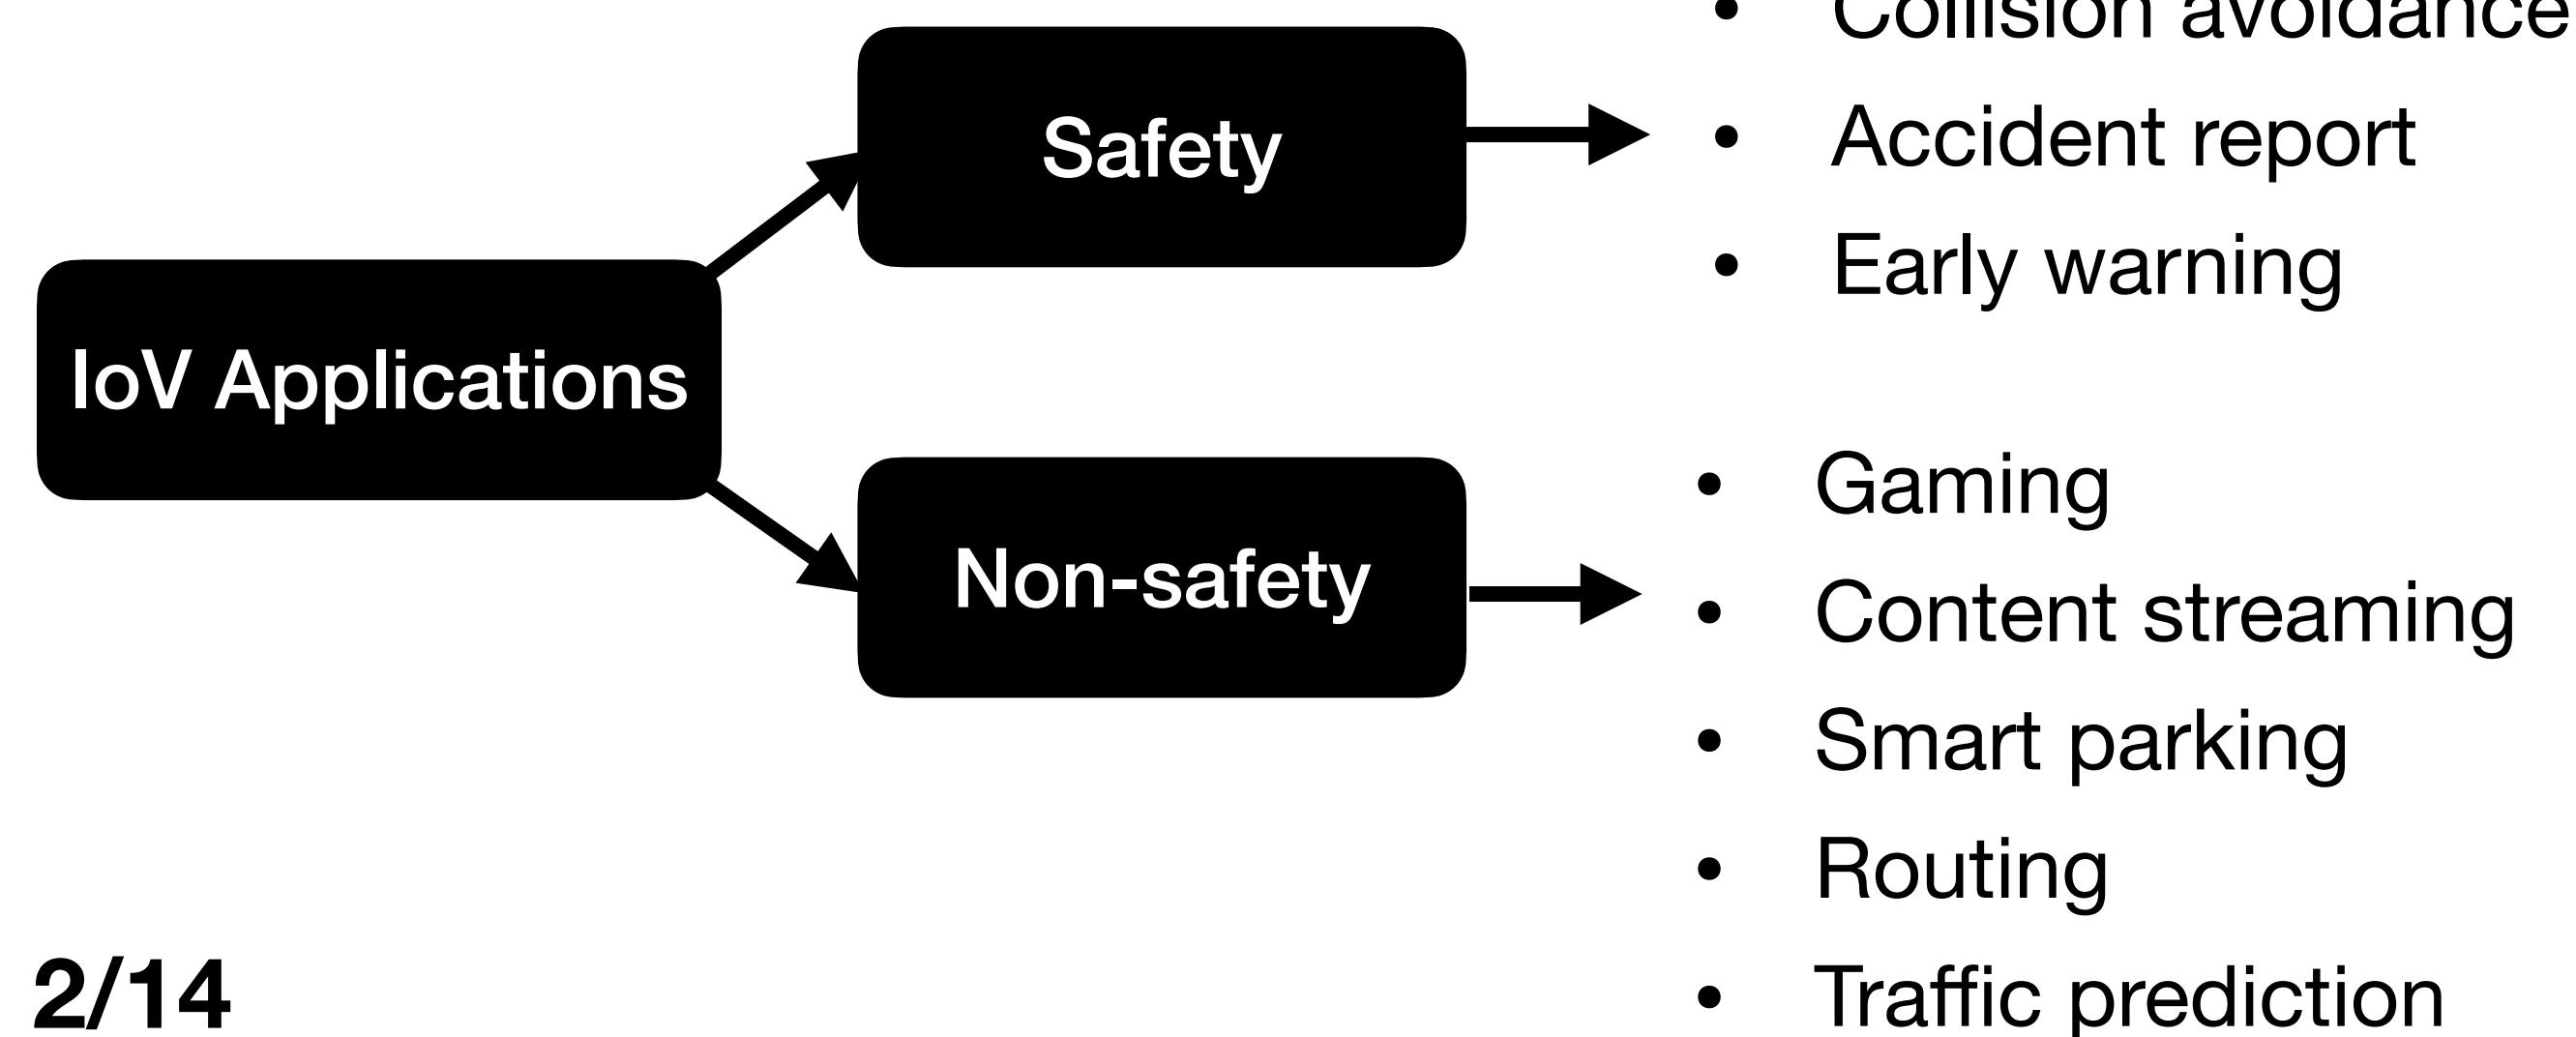
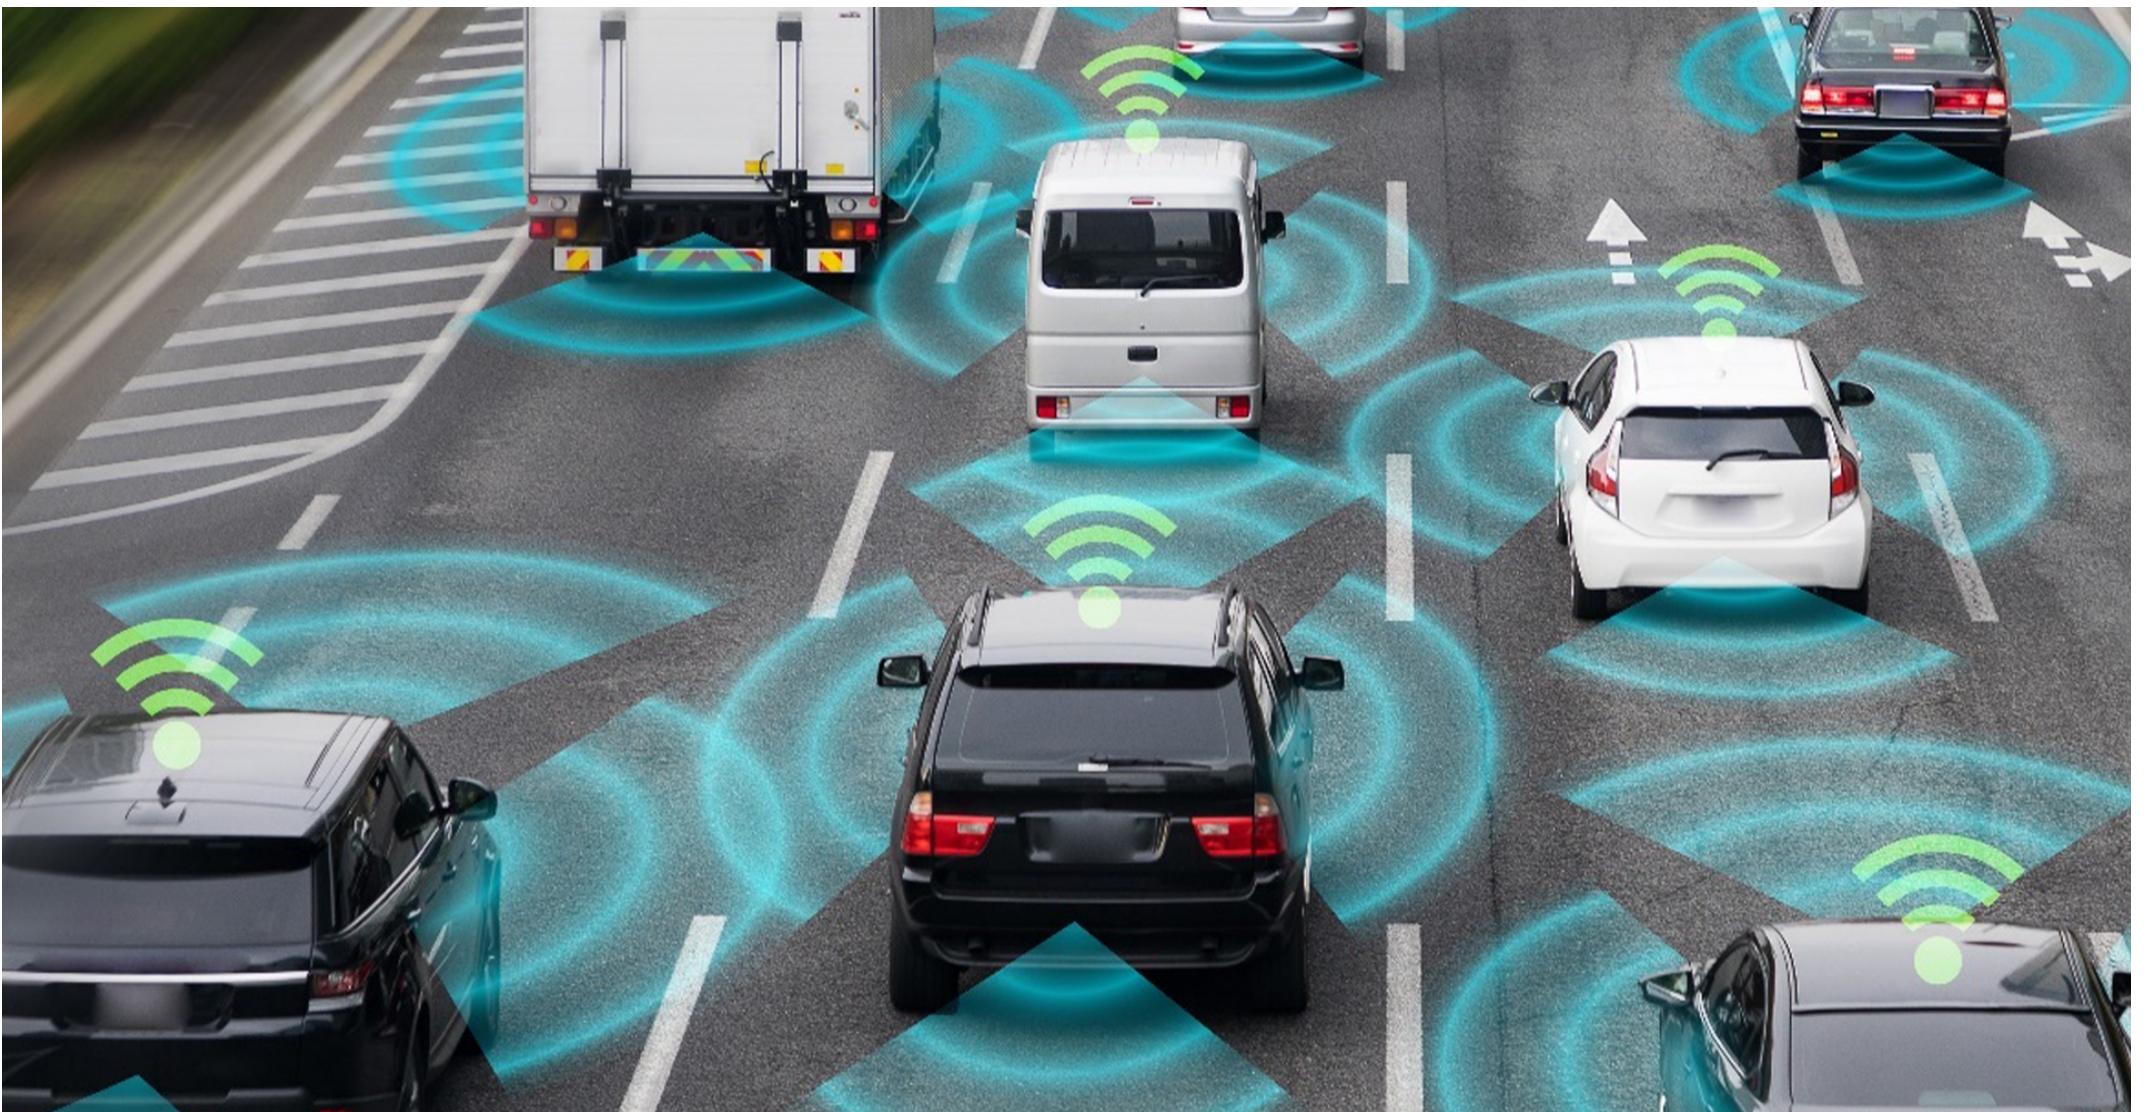
- Simulation Scenario
- Compared Algorithms
- Task Delay
- Queue Length

Part I: Introduction

The Internet of Vehicles

Why future cars will be connected

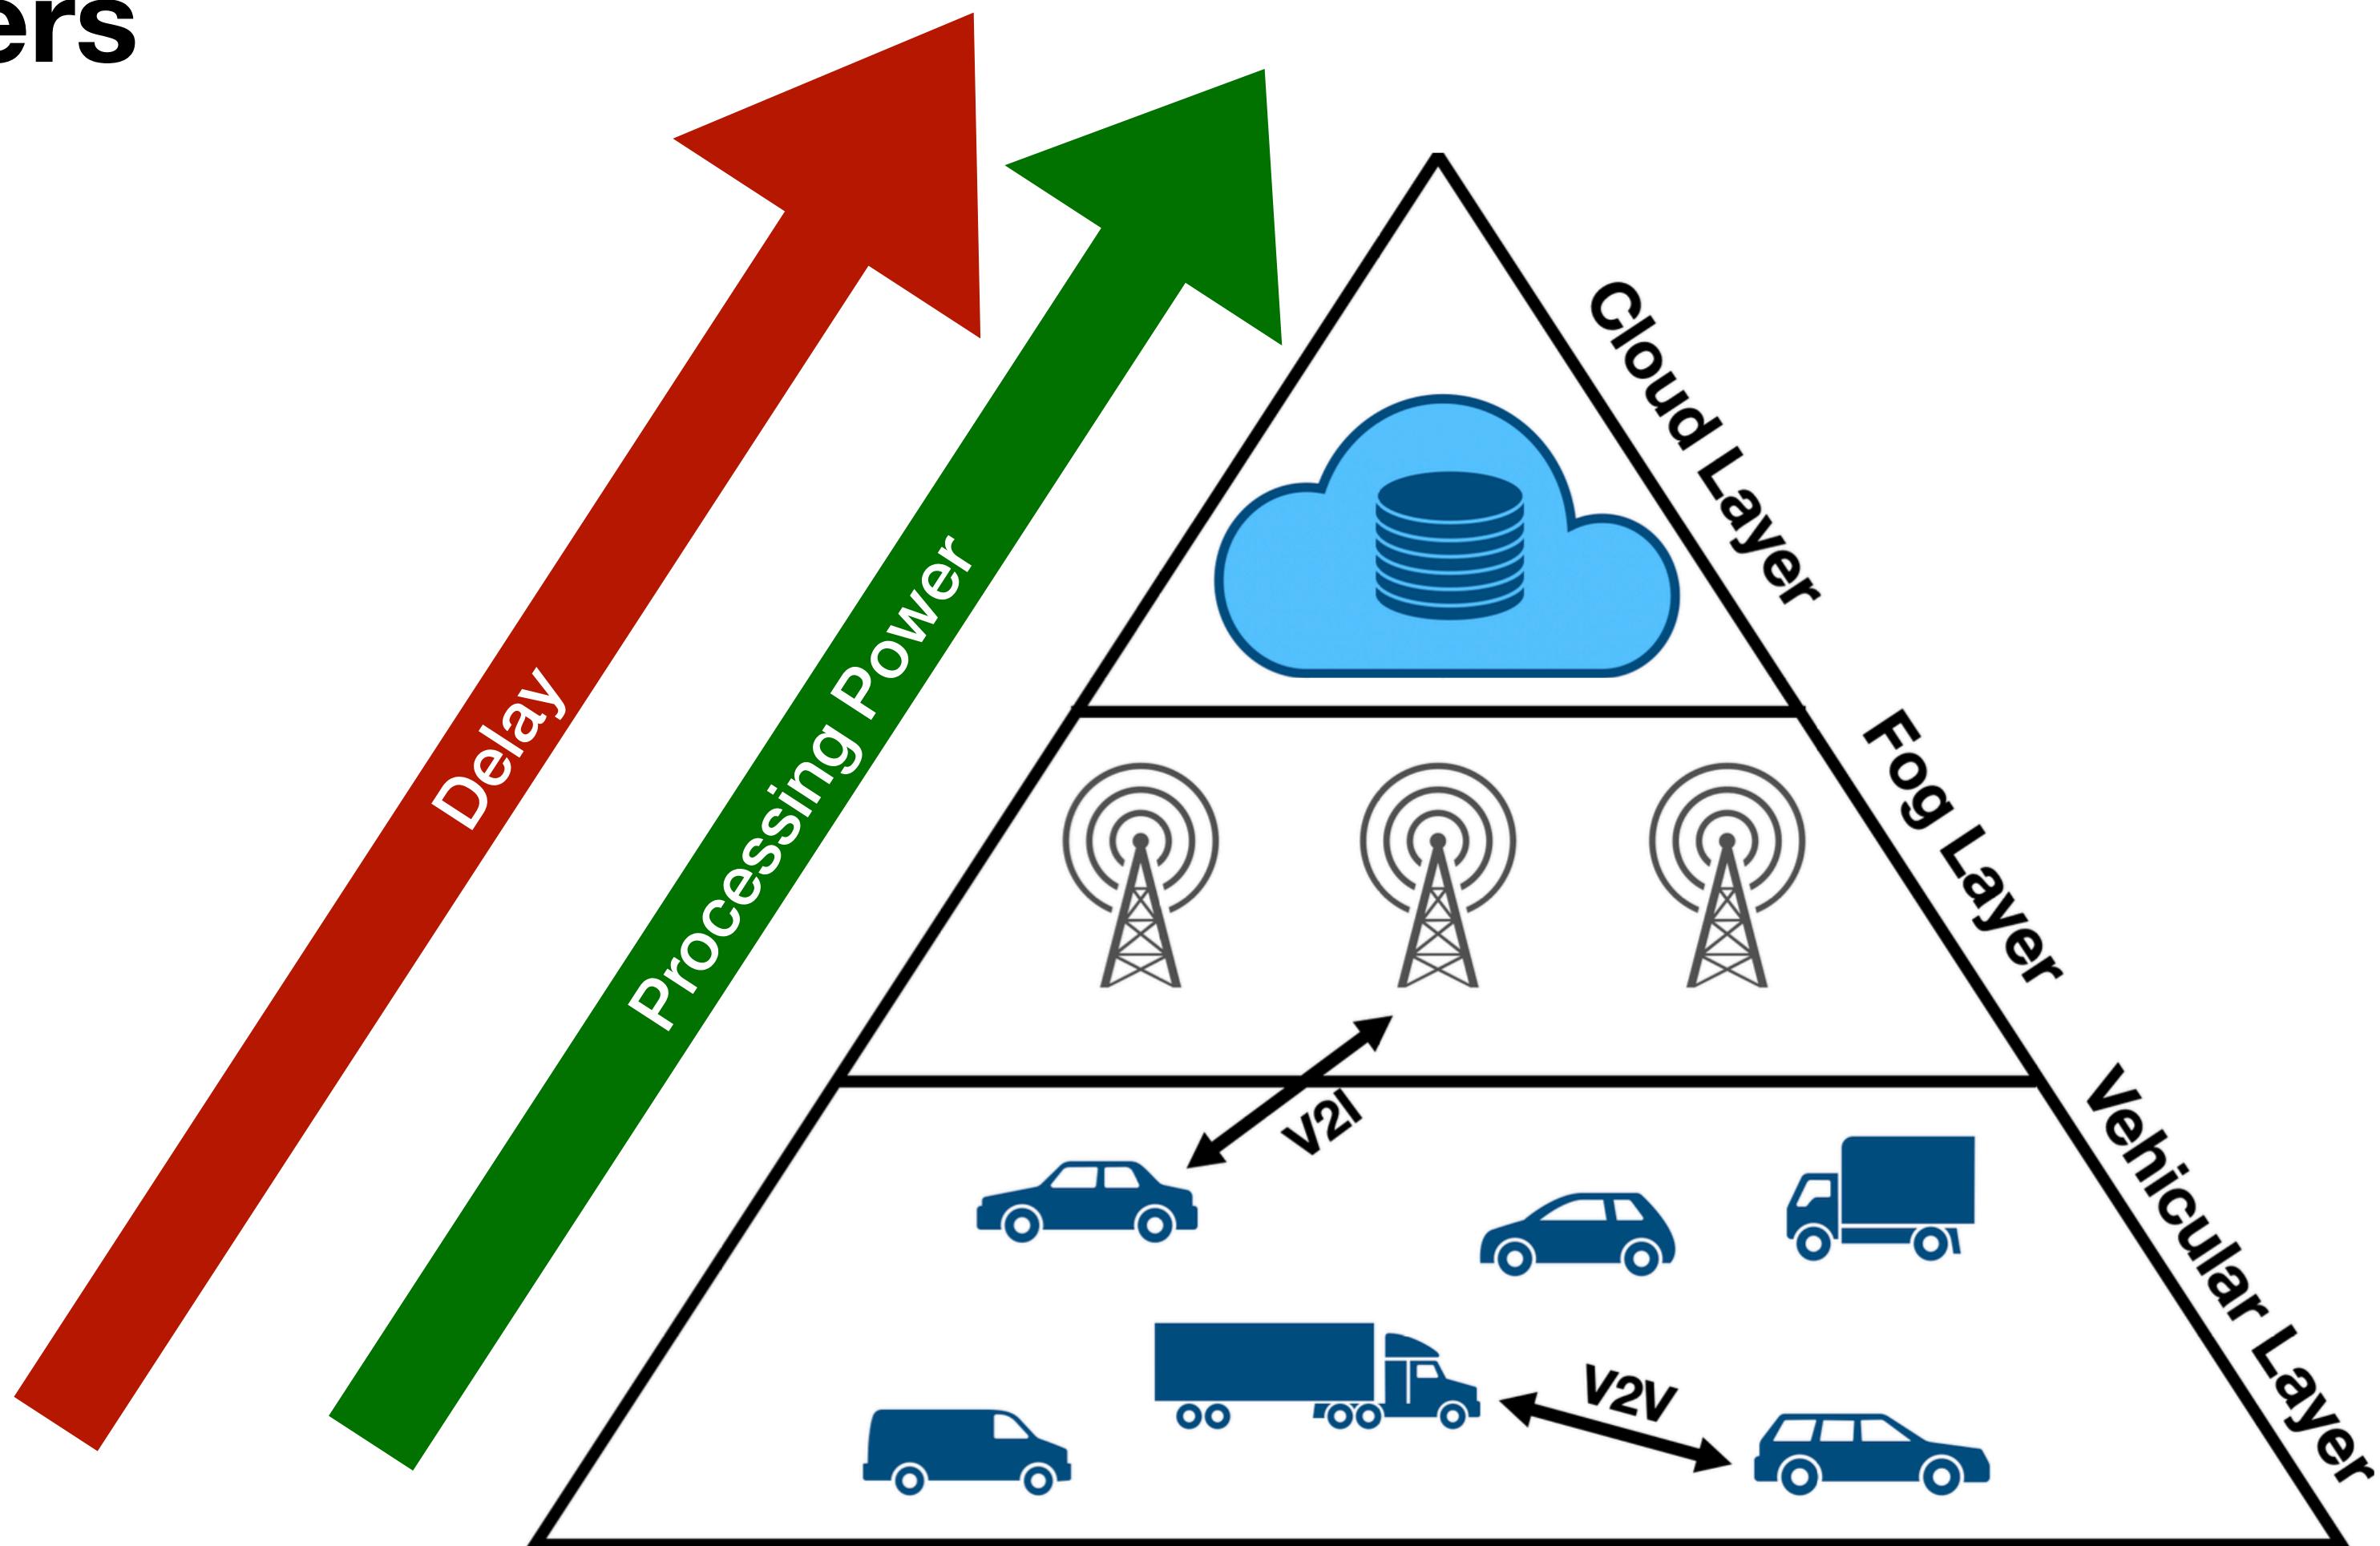
- 30 years ago, the main edge-devices were normal computers
- Now, almost everything is being connected, including **cars**
- Cars will communicate with each other, transmission towers, and the cloud, forming the **Internet of Vehicles (IoV)**



The Problem of Task Offloading

Offloading at Different Layers

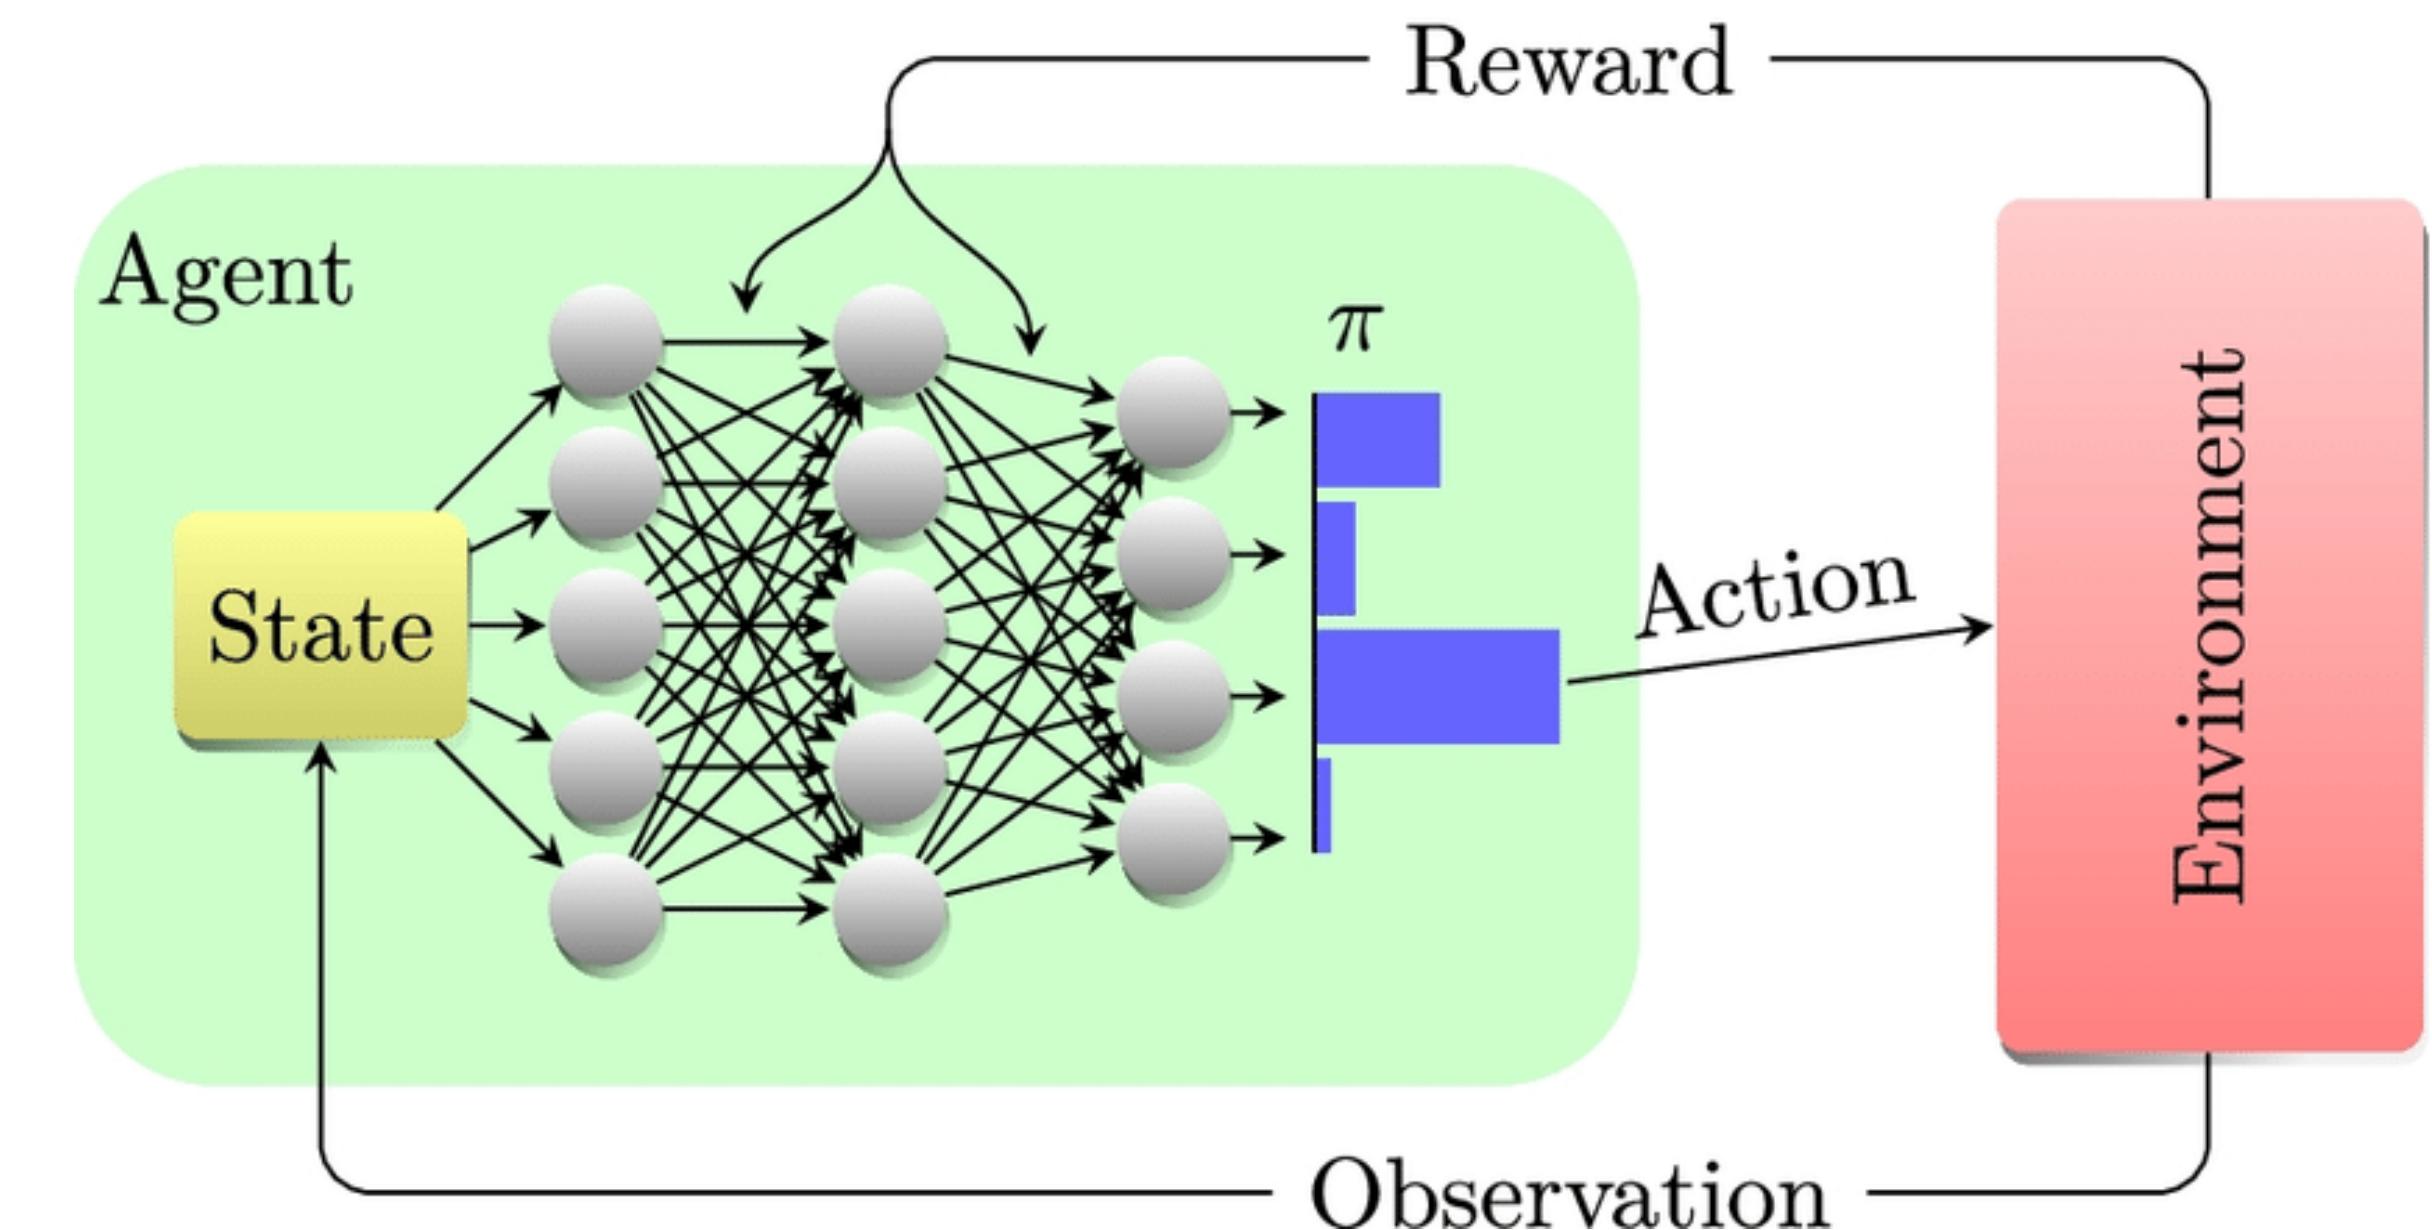
- So, future cars will need to do a lot of processing, but they can benefit from offloading
- **Task offloading:** sending a task to another layer for processing
 - **Edge:** the individual vehicles
 - **Fog:** the communication infrastructure
 - **Cloud:** the powerful servers at the network core
- **Optimization problem:** where to offload task?



Reinforcement Learning

The new paradigm of algorithm design

- **Old method:** humans explicitly think, implement and redesign algorithms
- **New method:** just specify the problem and let a computer discover near-optimal algorithms
- **Reinforcement Learning (RL):** subset of ML instructing agents on problem-solving through repeated trial and error
- **Markov Decision Process (MDP):** mathematical framework of RL



Part II: Problem Formulation

Problem Formulation

Vehicle model

- Vehicles are either moving or parked
- Vehicles are characterized by:
 - (X, Y) coordinates
 - Velocity
 - CPU processing power (based on real hardware)
- Vehicles move randomly in a grid-planned city (common in the USA)

TABLE III: Processor Specifications

Processor Name	Processor Speed (MIPS)
ARM Cortex A72	18,375
AMD Phenom II X4 940	42,820
ARM Cortex A73	71,120

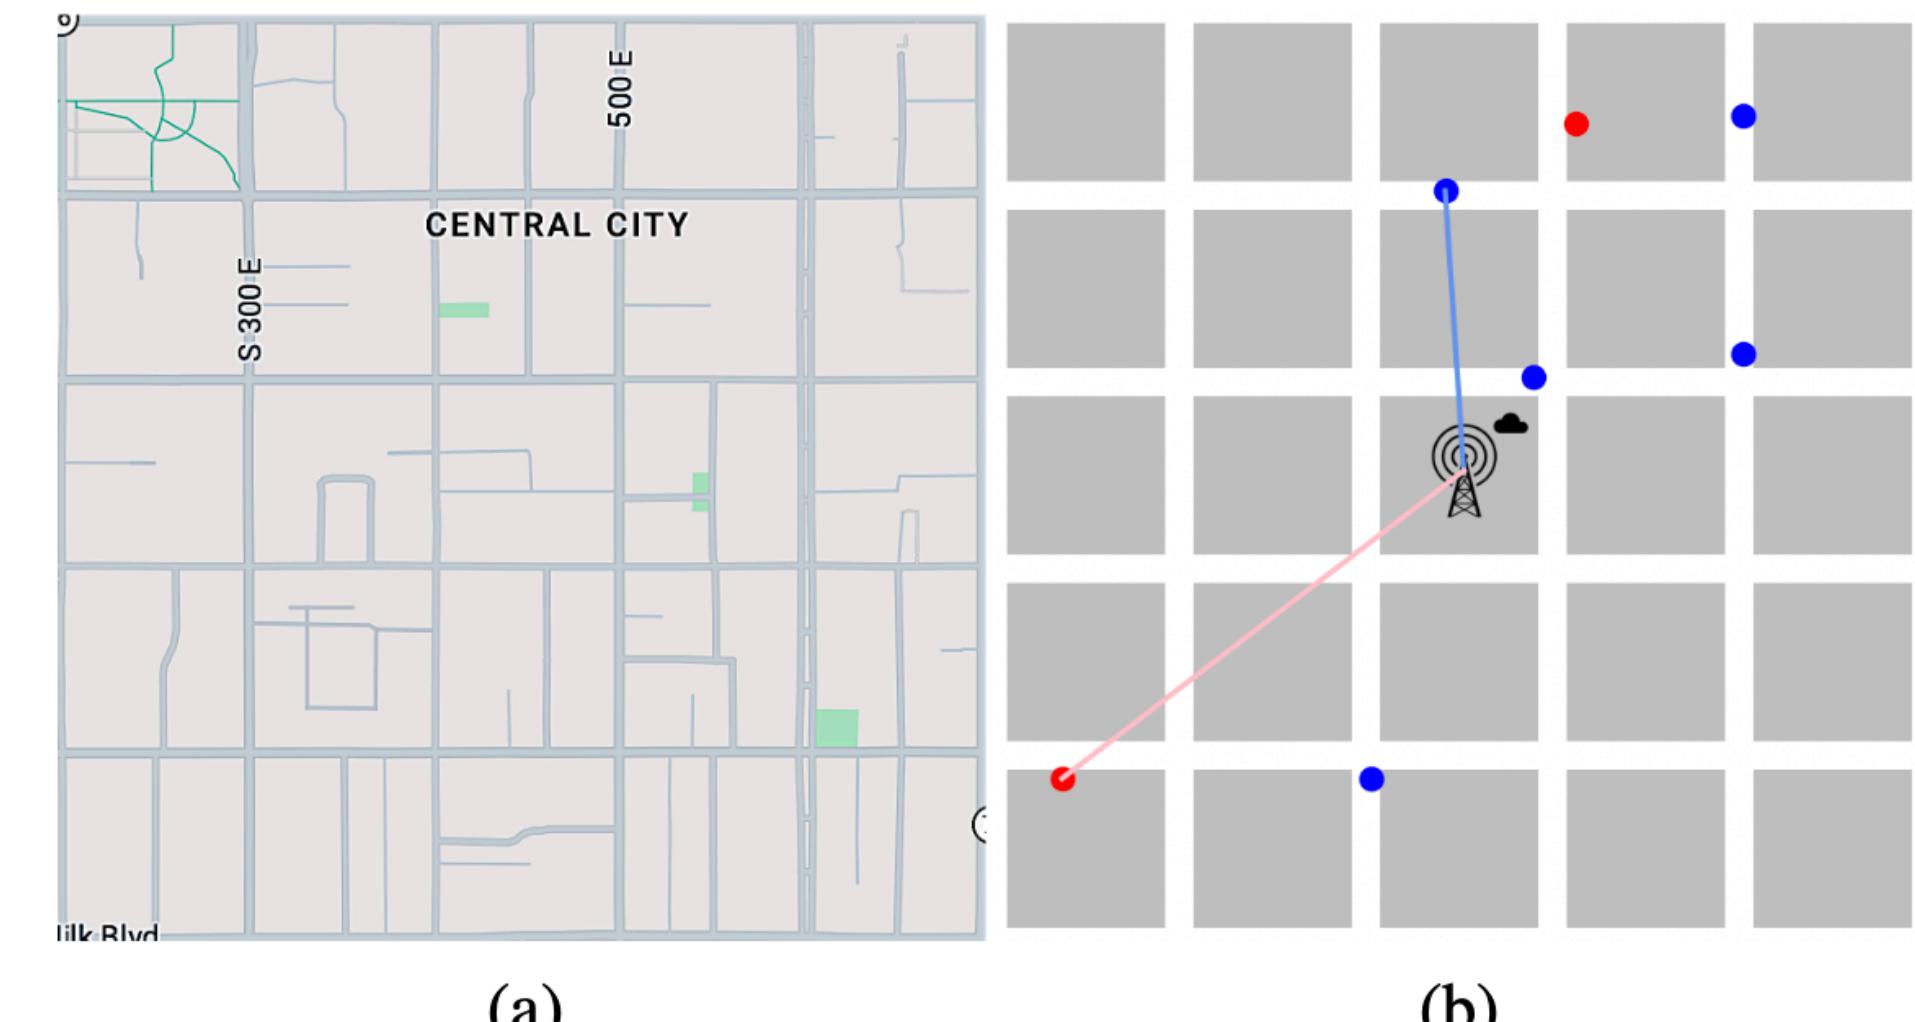


Fig. 3: Downtown of Salt Lake City, U.S., showcasing its grid-like structure of 200 by 200 meter city blocks (a) (Map data @2024 Google). Simulated environment based on Salt Lake City blocks (b).

Problem Formulation (Cont.)

Task model

- **Tasks:** a well-defined set of operations which a vehicle creates
- Tasks are characterized by:
 - Instruction count (MI)
 - Size (Mb)
- Tasks are based on **SPEC CPU Dataset**, consisting of 18 tasks with varying computational complexities

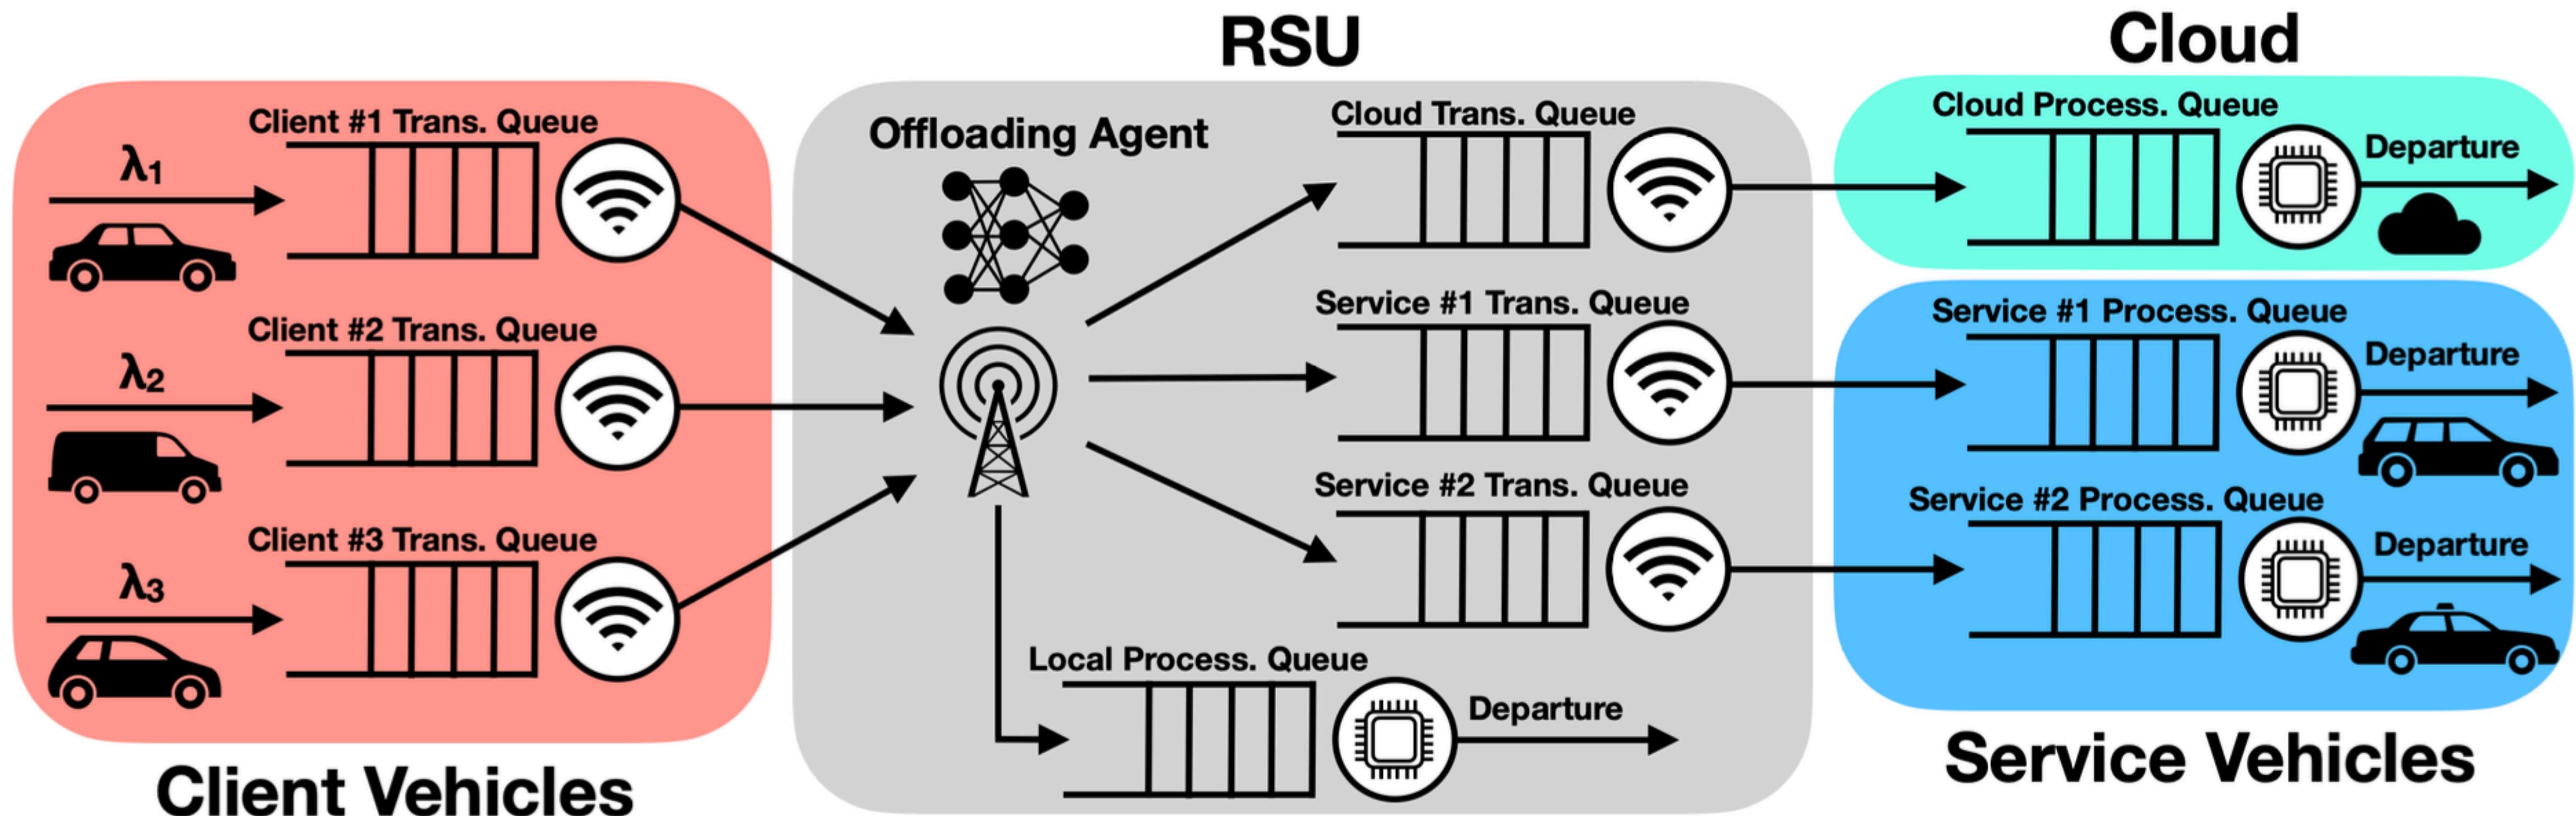
SPEC CPU95			
Program	Input	INT/ FP	Dynamic Instruction Count
li	*.lsp	INT	75.6 billion
m88ksim	ctl.in	INT	520.4 billion
compress	bigtest.in	INT	69.3 billion
ijpeg	penguin.ppm	INT	41.4 billion
gcc	expr.i	INT	1.1 billion
perl	perl.in	INT	16.8 billion
vortex	*	INT	*
wave5	wave5.in	FP	30 billion
hydro2d	Hydro2d.in	FP	44 billion
swim	swim.in	FP	30.1 billion
applu	Applu.in	FP	43.7 billion
mgrid	Mgrid.in	FP	56.4 billion
turb3d	turb3d.in	FP	91.9
su2cor	su2cor.in	FP	33 billion
fpppp	natmos.in	FP	116 billion
apsi	apsi.in	FP	28.9 billion
tomcatv	tomcatv.in	FP	26.3 billion

Problem Formulation (Cont.)

Communication model

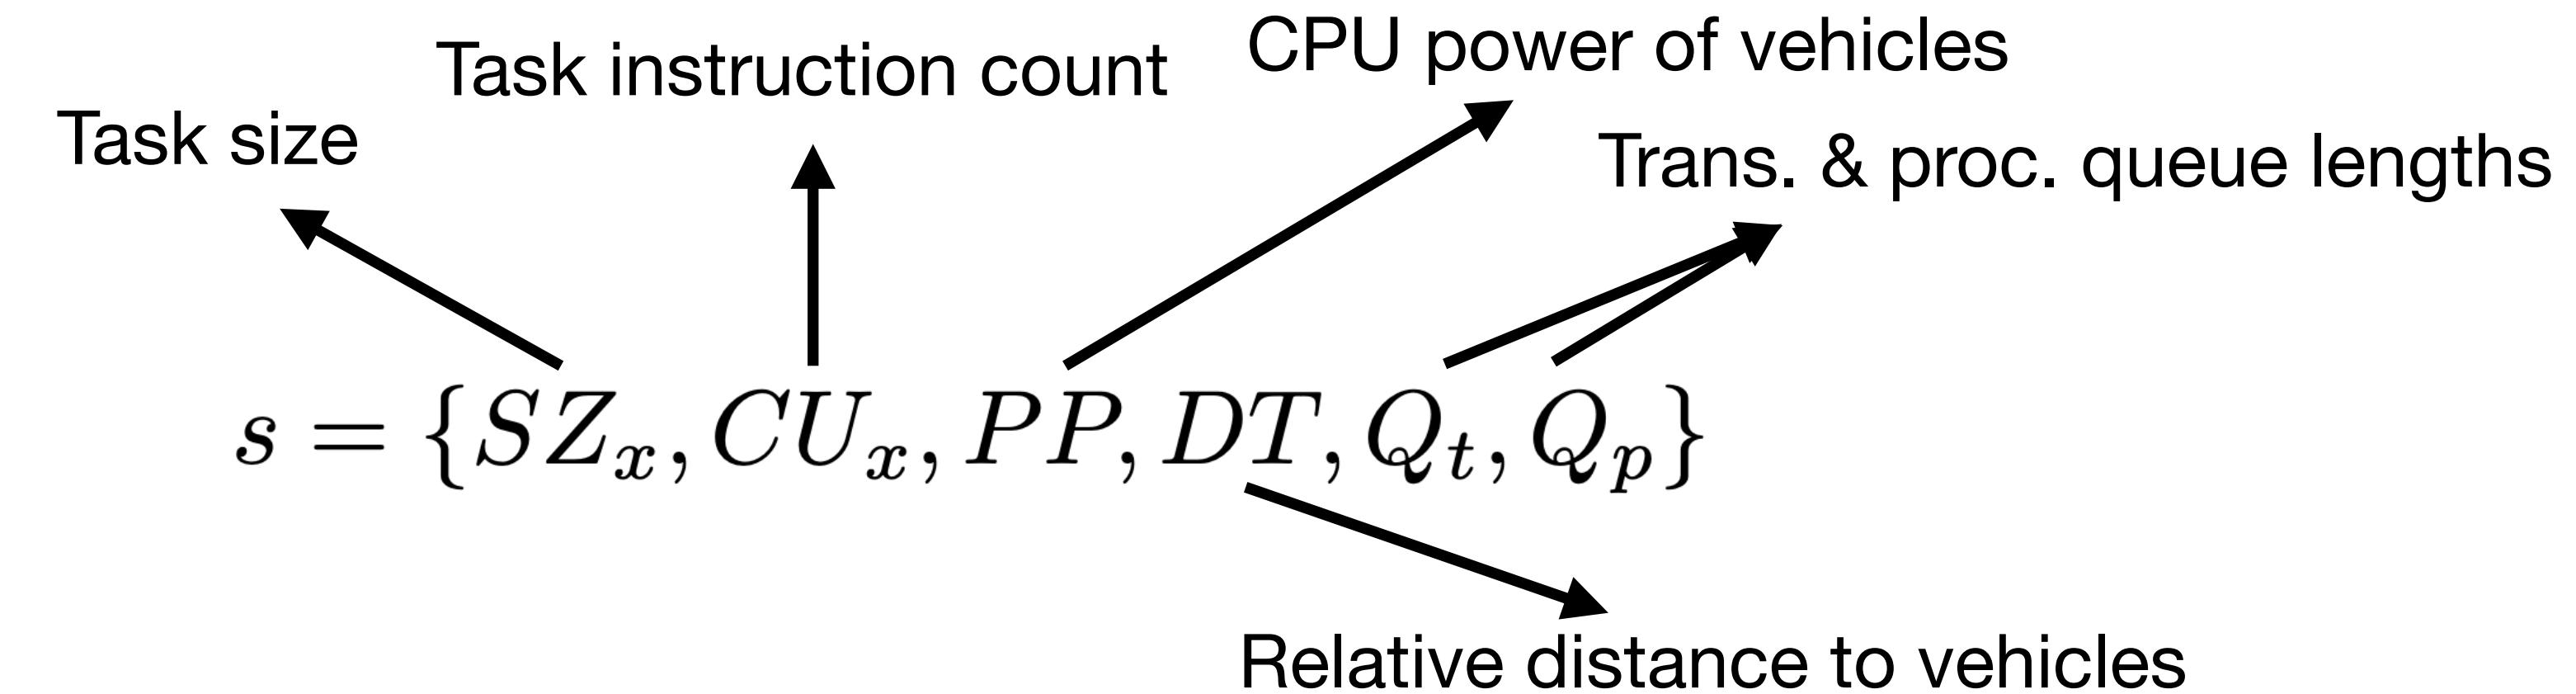
- The total delay experienced by a task is the sum of the following:

$$d_{\text{total}} = d_{\text{client}} + d_{\text{rsu}} + d_{\text{service}}$$



MDP Specification

- State space:



- Action space:

$$A = V \cup \{a_{\text{cloud}}\} \cup \{a_{\text{rsu}}\}$$

- Reward function

$$R(s, a) = \sum_{x \in U} -d_{\text{total}}^x$$

↑
Negative of the sum of all task delays

Part III: Results

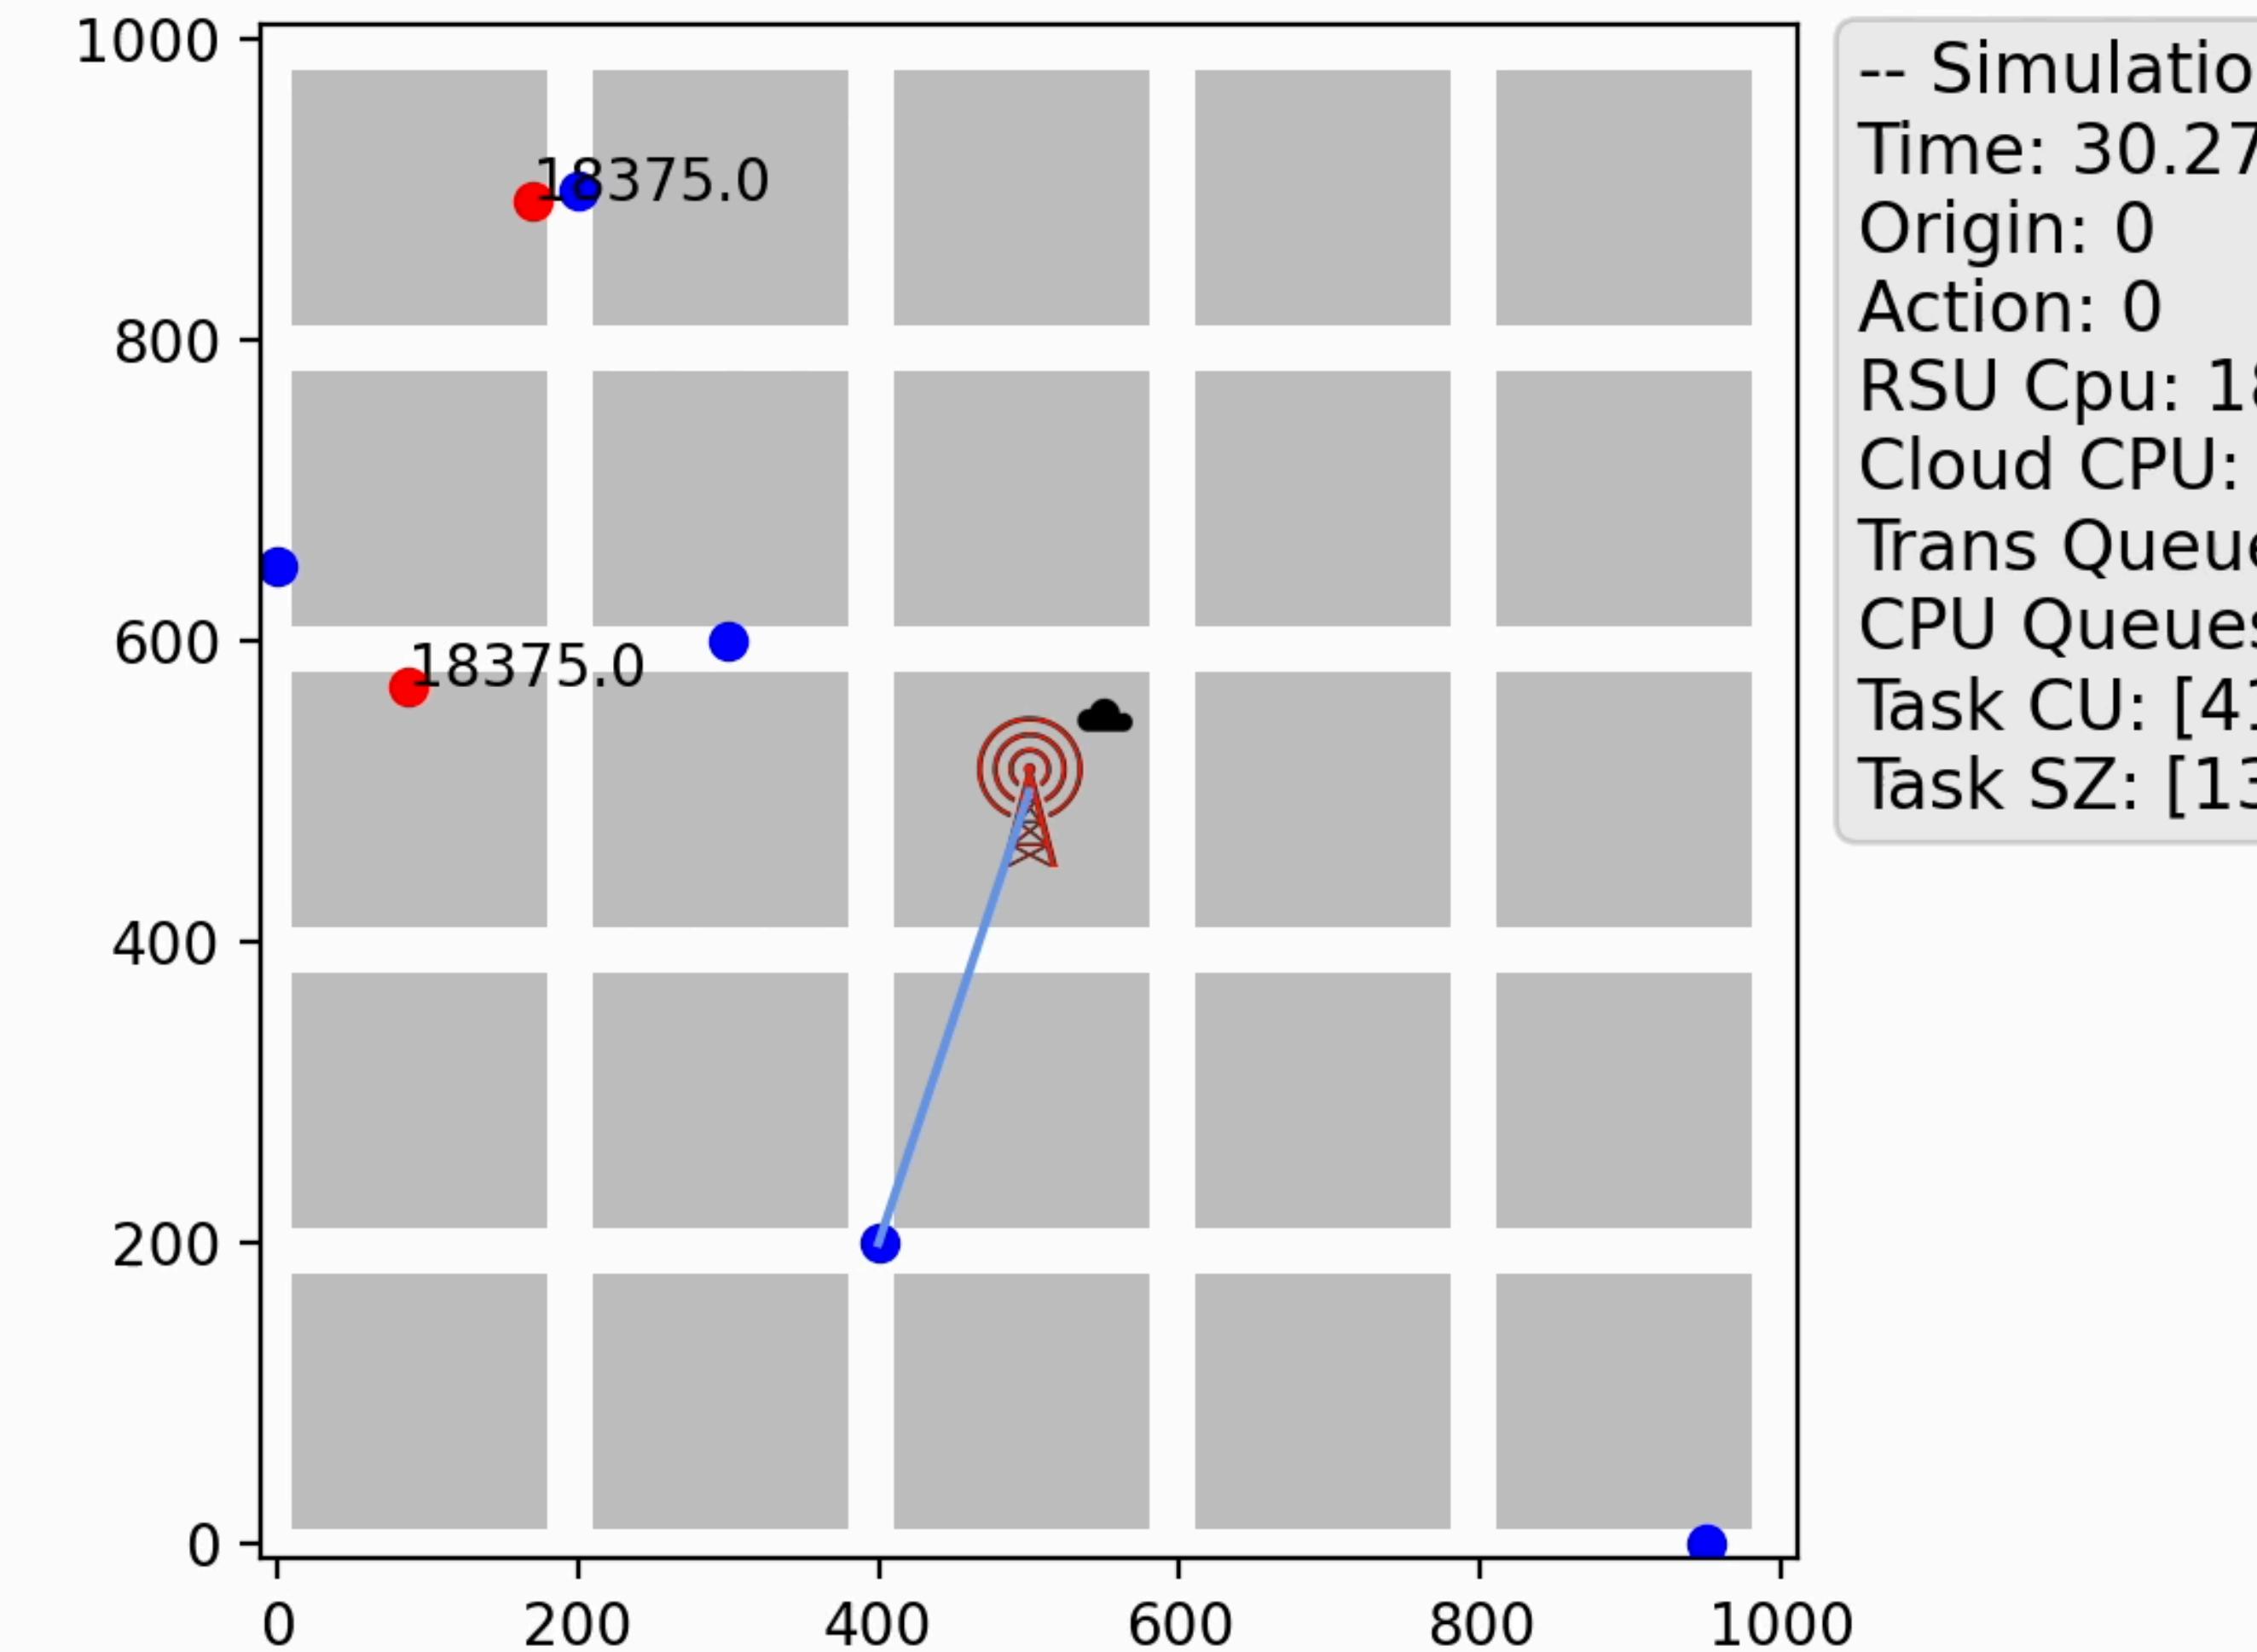
Simulation Scenarios

- Three different scenarios considered:
 - Scenario 1: 5 moving vehicles & 2 parked vehicles
 - Scenario 2: 10 moving vehicles & 4 parked vehicles
 - Scenario 3: 20 moving vehicles & 8 parked vehicles
- Four algorithms compared:
 - Deep RL (PPO)
 - Random
 - Greedy
 - Cloud-only

TABLE II: Environment Settings

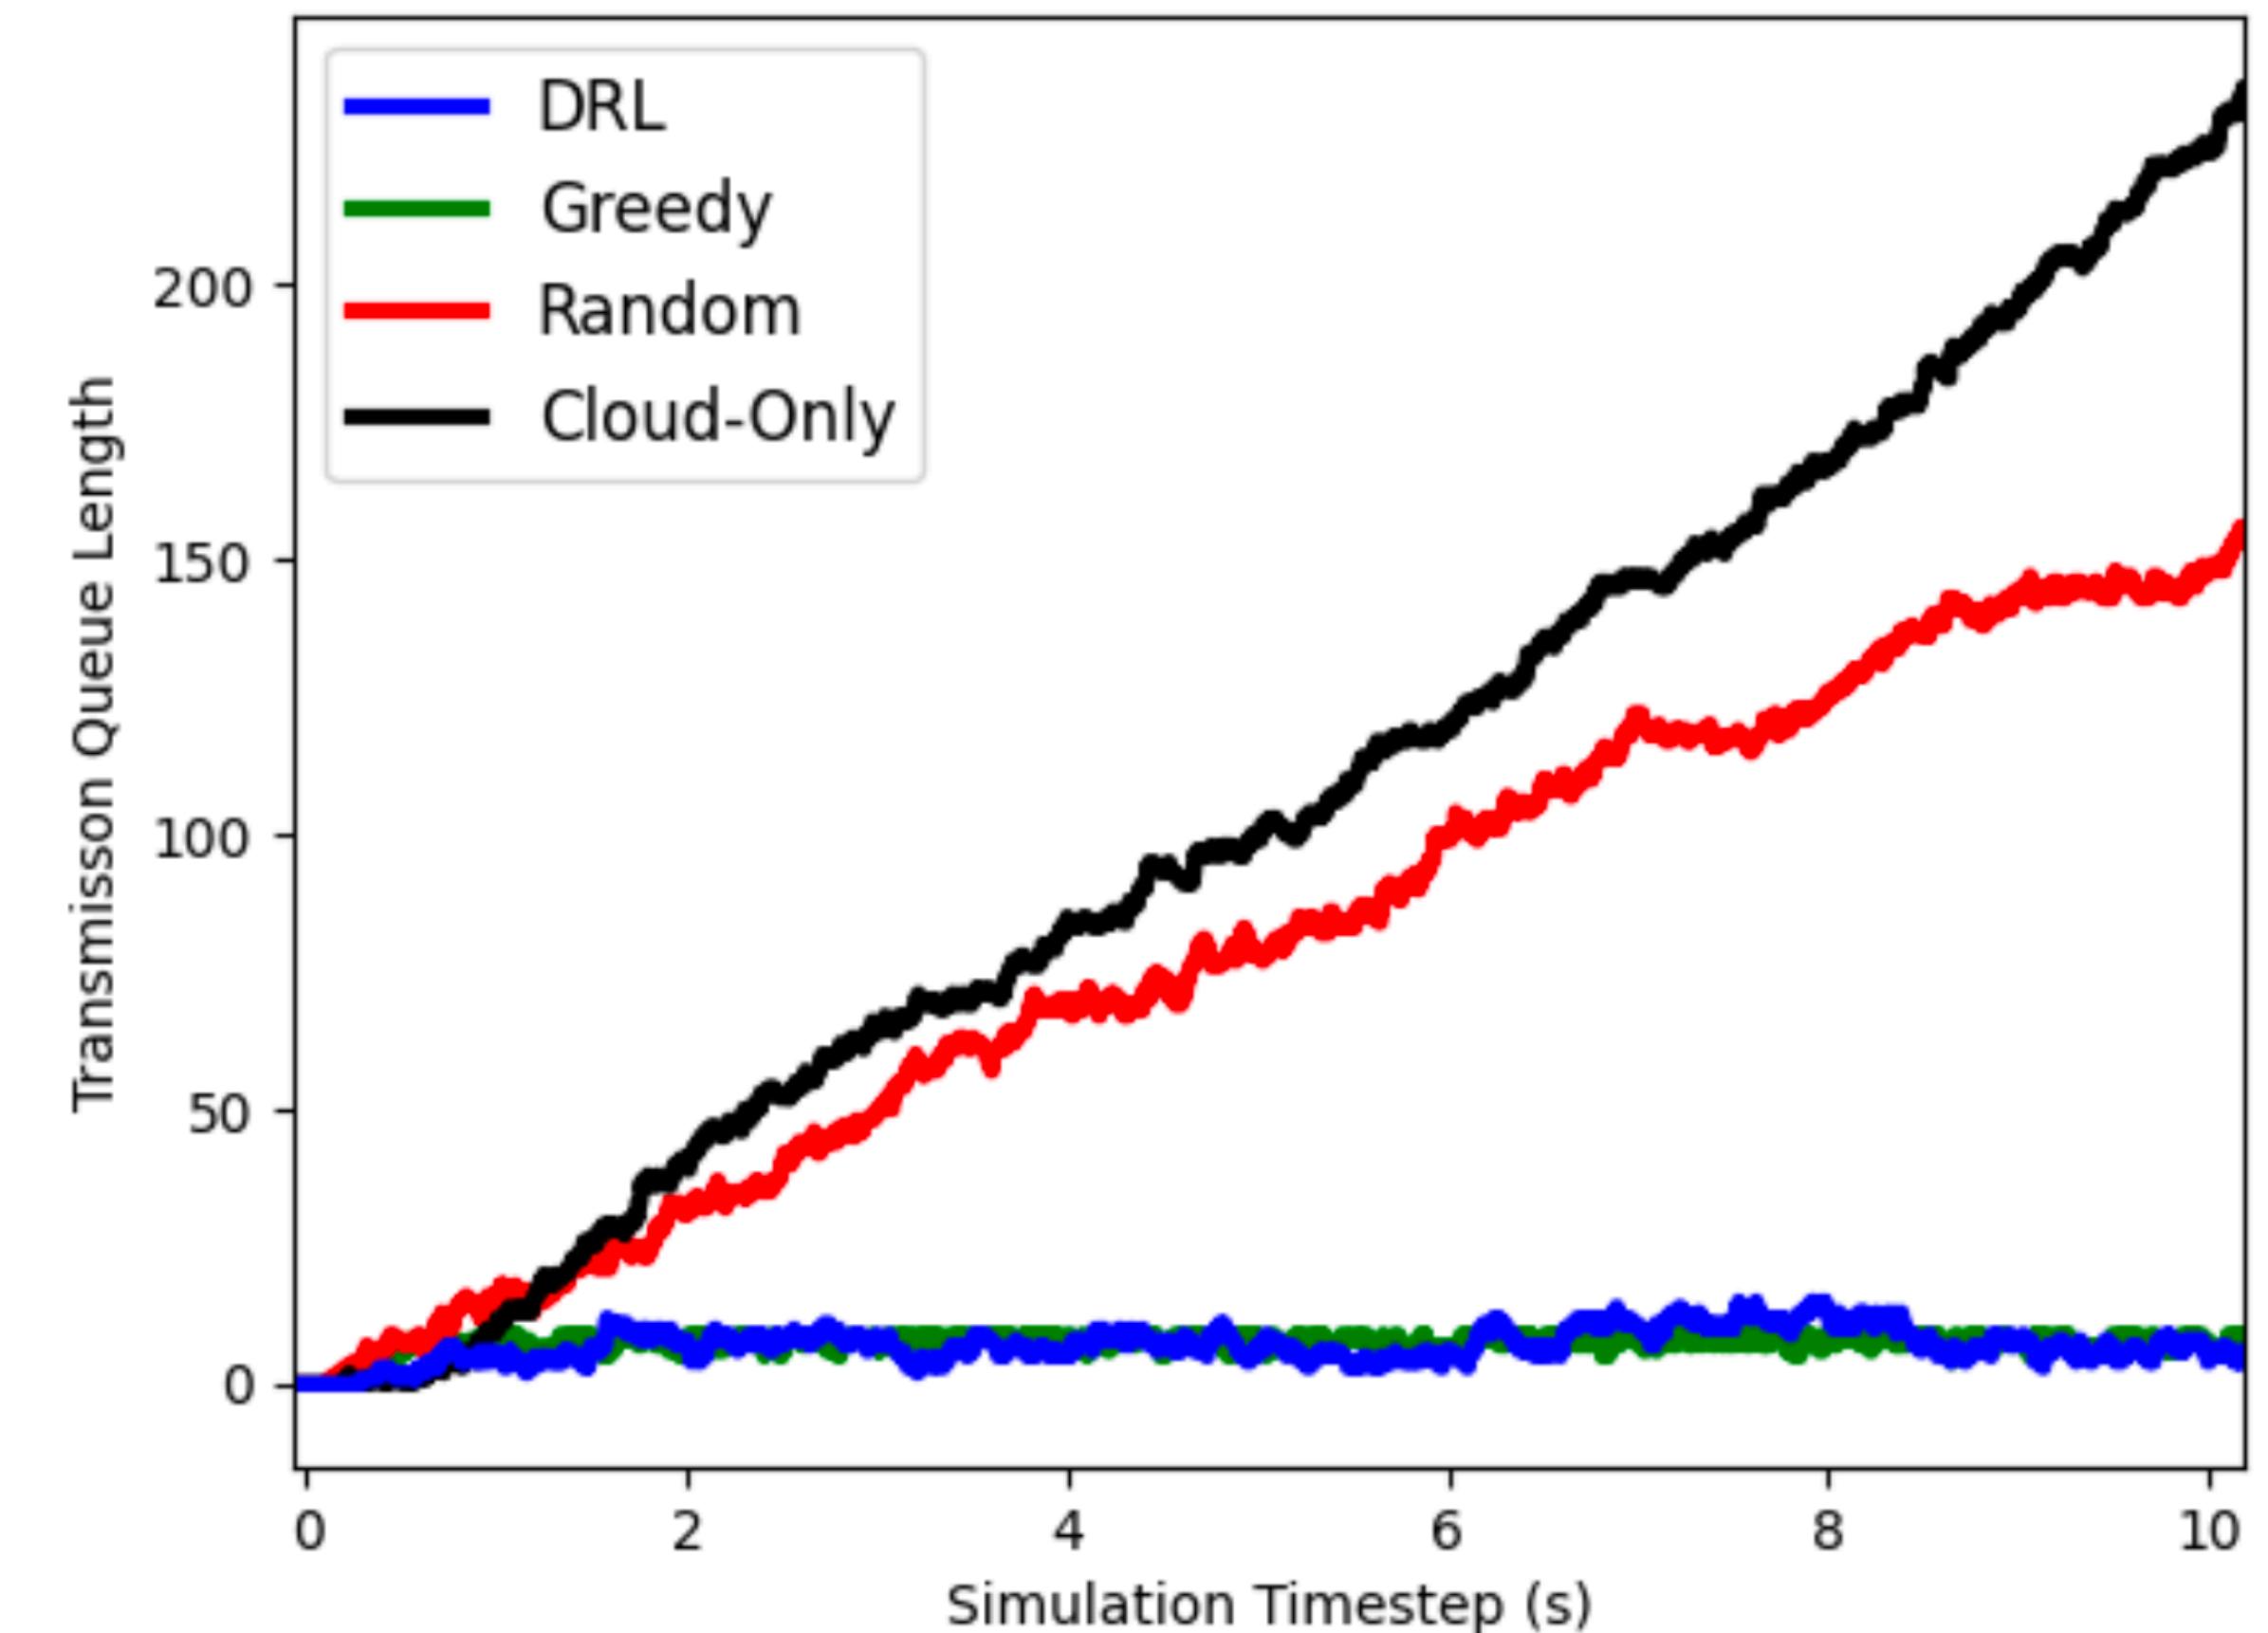
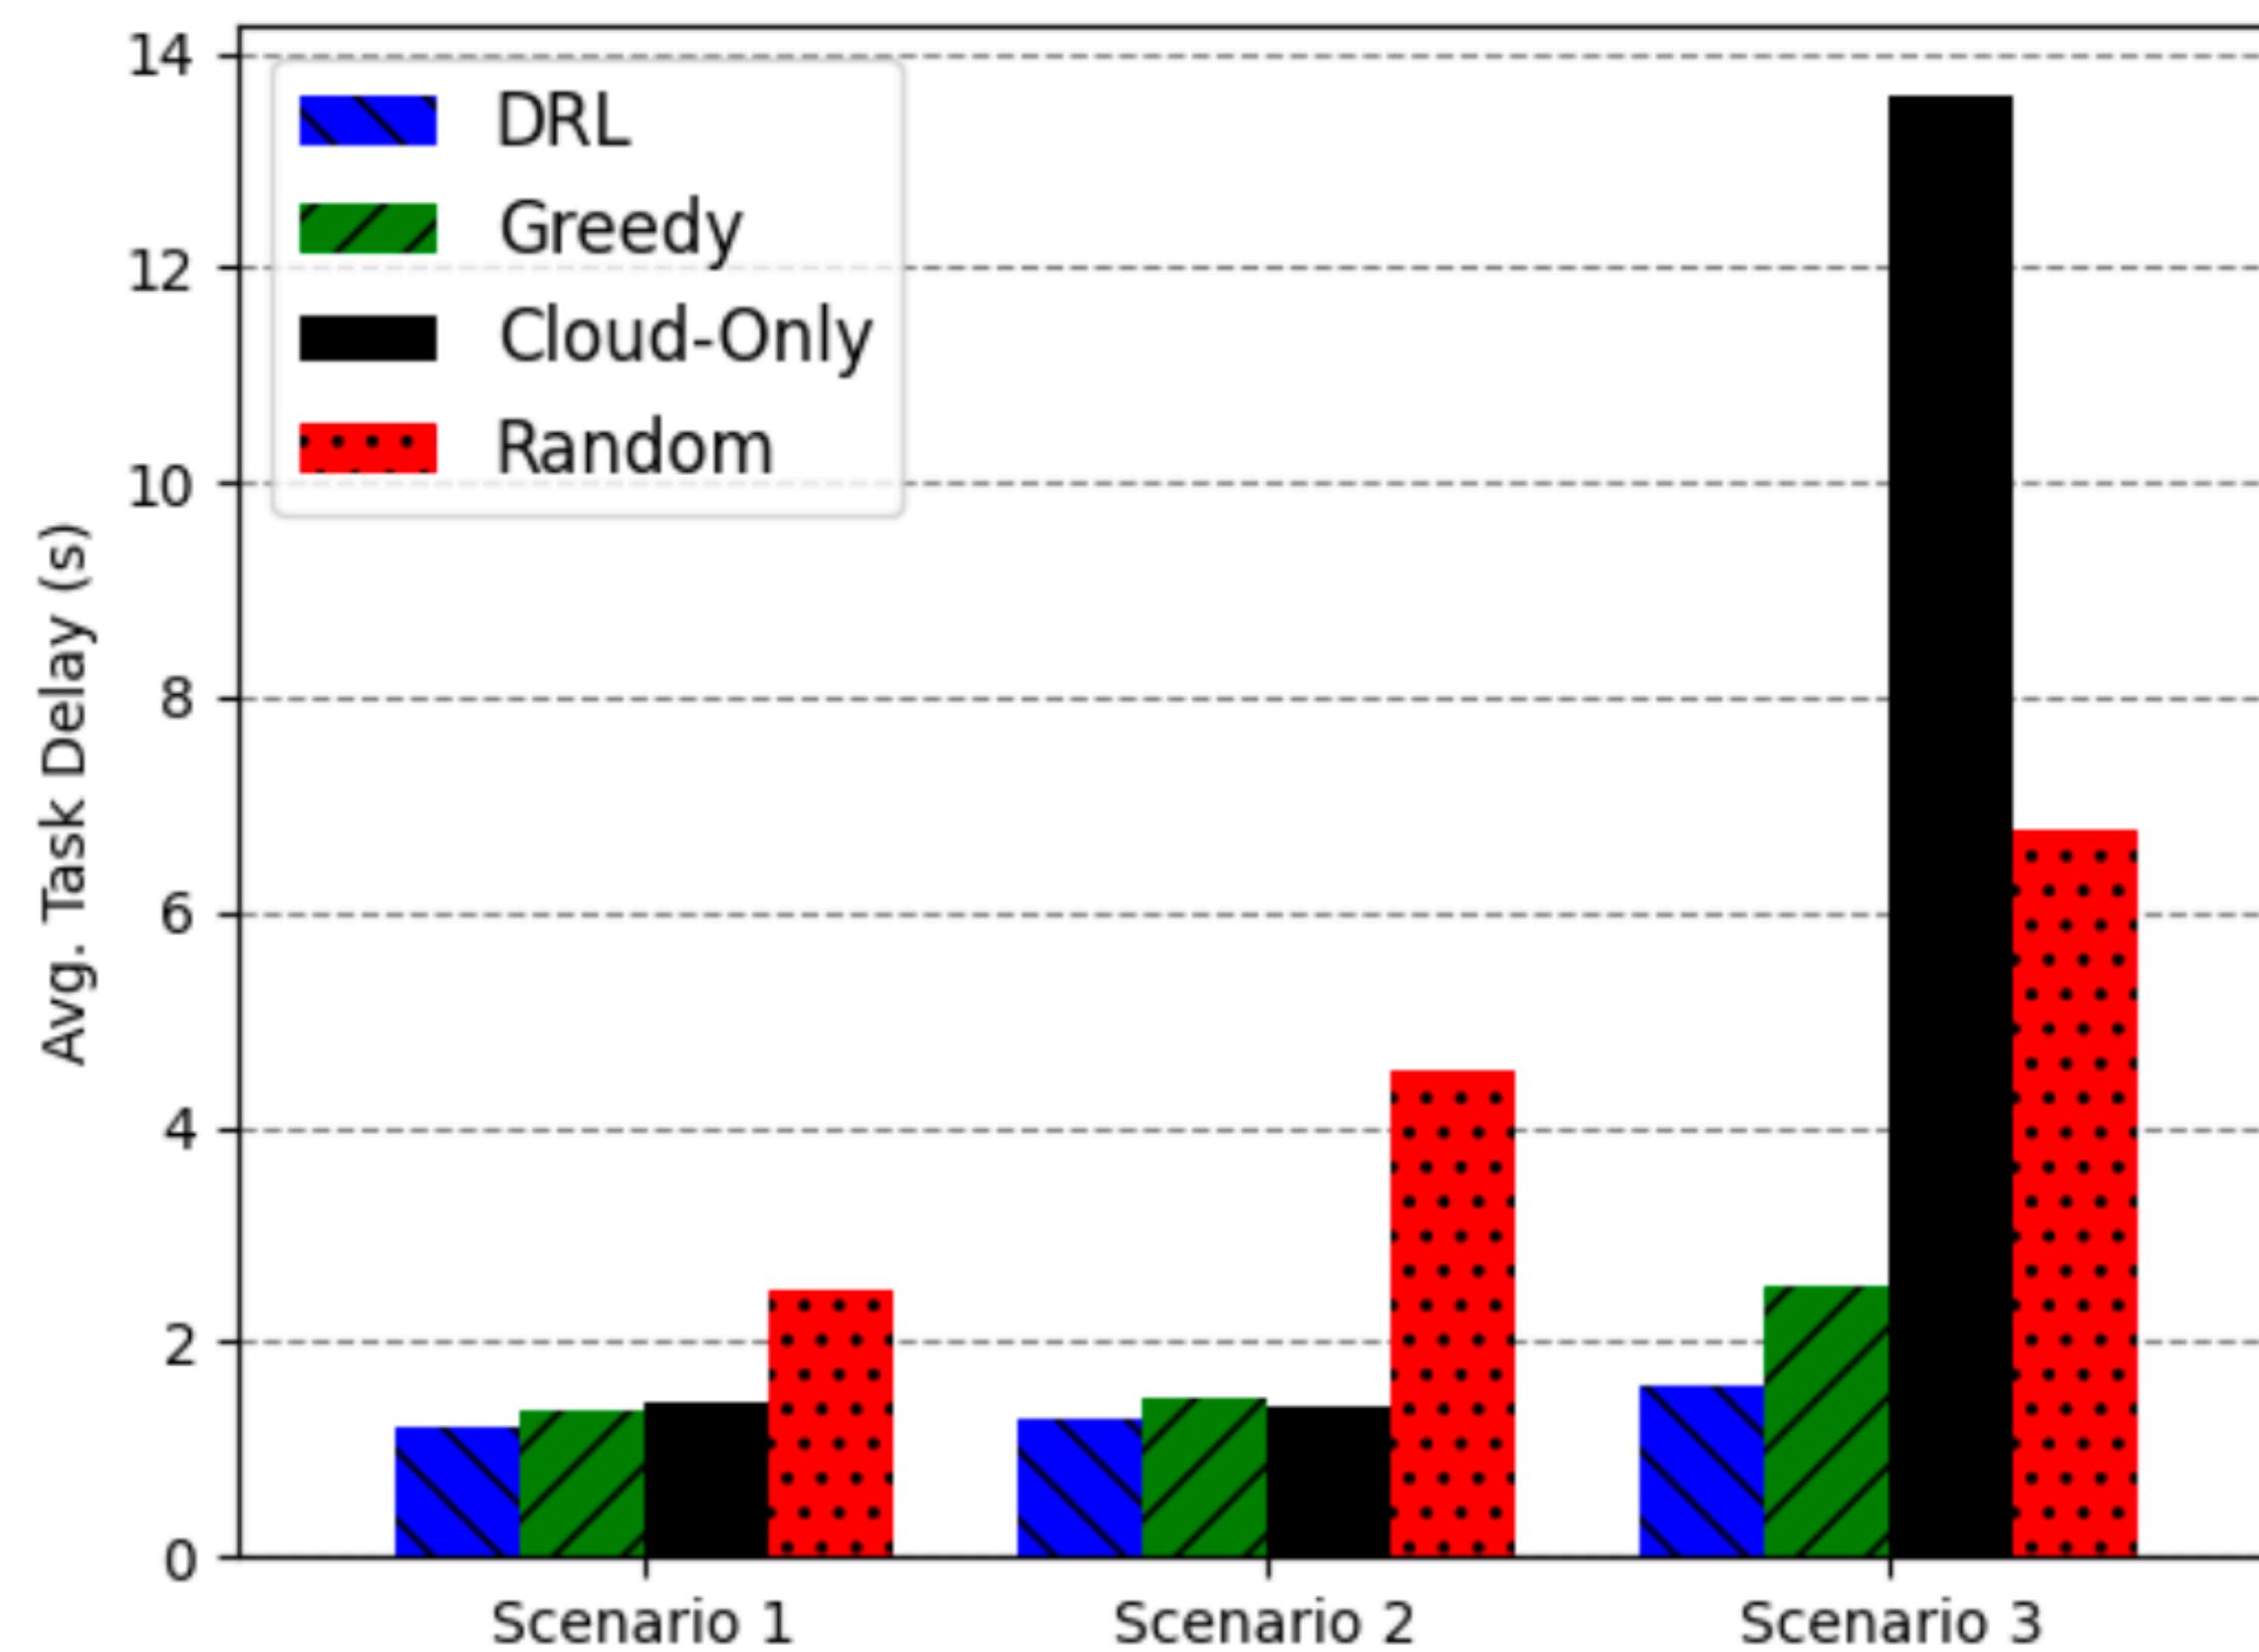
Parameter	Value
$l \times w$	1000 (m) \times 1000 (m)
n	5, 10, 20
m	2, 4, 8
λ_i	3 (tasks per sec)
vel_i	$\sim U\{10, 20, 25, 40\}$ (km/h)
SZ_x	$\sim U\{20, 40\}$ (Mb)
B	40 (MHz)
P_T	1 (W)
G_T, G_R	5 (dBi)
N_0	174 (dBm/Hz)
CPU_{rsu}	$1 \times 18,375$ (MIPS)
CPU_{cloud}	$\infty \times 100,000$ (MIPS)

Simulation Animation



Results

Average task delay & queue length



Code Repository

Open-source, reproducible research

The screenshot shows the GitHub repository page for "VFC-Offloading-RL". The repository is public and has 10 stars, 2 forks, and 14 commits. The main branch is "main". The repository contains files like README.md, .gitignore, requirements.txt, and system-model-visual.png. A detailed description of the project is provided, mentioning "Deep Reinforcement Learning for Delay-Optimized Task Offloading in Vehicular Fog Computing". A diagram at the bottom illustrates the architecture of the Vehicular Fog Computing Task Offloading RL Environment.

Vehicular Fog Computing Task Offloading RL Environment

This repository contains the source code of the following paper: "Deep Reinforcement Learning for Delay-Optimized Task Offloading in Vehicular Fog Computing". The pre-print version can be view here: <https://arxiv.org/abs/2410.03472>. We appreciate citing the paper if you found this repository to be useful.

The diagram illustrates the architecture of the Vehicular Fog Computing Task Offloading RL Environment. It shows three main components: Client Vehicles, RSU (Road Side Unit), and Service Vehicles. Client Vehicles (Client #1, Client #2, Client #3) send tasks with rates λ_1 , λ_2 , and λ_3 to the RSU. The RSU contains an Offloading Agent and a Local Process. Queue. Tasks from Client Vehicles enter the RSU via WiFi. The RSU then processes tasks and sends them to the Cloud or Service Vehicles via WiFi. The Cloud contains a Cloud Trans. Queue and a Cloud Process. Queue, leading to Departure. Service Vehicles contain Service Trans. Queues and Process. Queues, also leading to Departure. The RSU also has a Local Process. Queue leading to Departure.

Github repository link:
github.com/Procedurally-Generated-Human/VFC-Offloading-RL

Current stats:
11 stars
2 forks

The End
Thank You

Extra Slide #1

Delay model

$$d_{\text{total}} = d_{\text{client}} + d_{\text{rsu}} + d_{\text{service}}$$

$$d_{\text{client}} = d_{\text{client-queue}} + d_{\text{client-trans}}$$

$$d_{\text{rsu}} = d_{\text{rsu-queue}} + d_{\text{rsu-trans}}$$

$$d_{\text{service}} = d_{\text{service-queue}} + d_{\text{service-process}}$$

$$d_{\text{service-process}} = \frac{CU_x}{CPU_{\text{service}}}$$

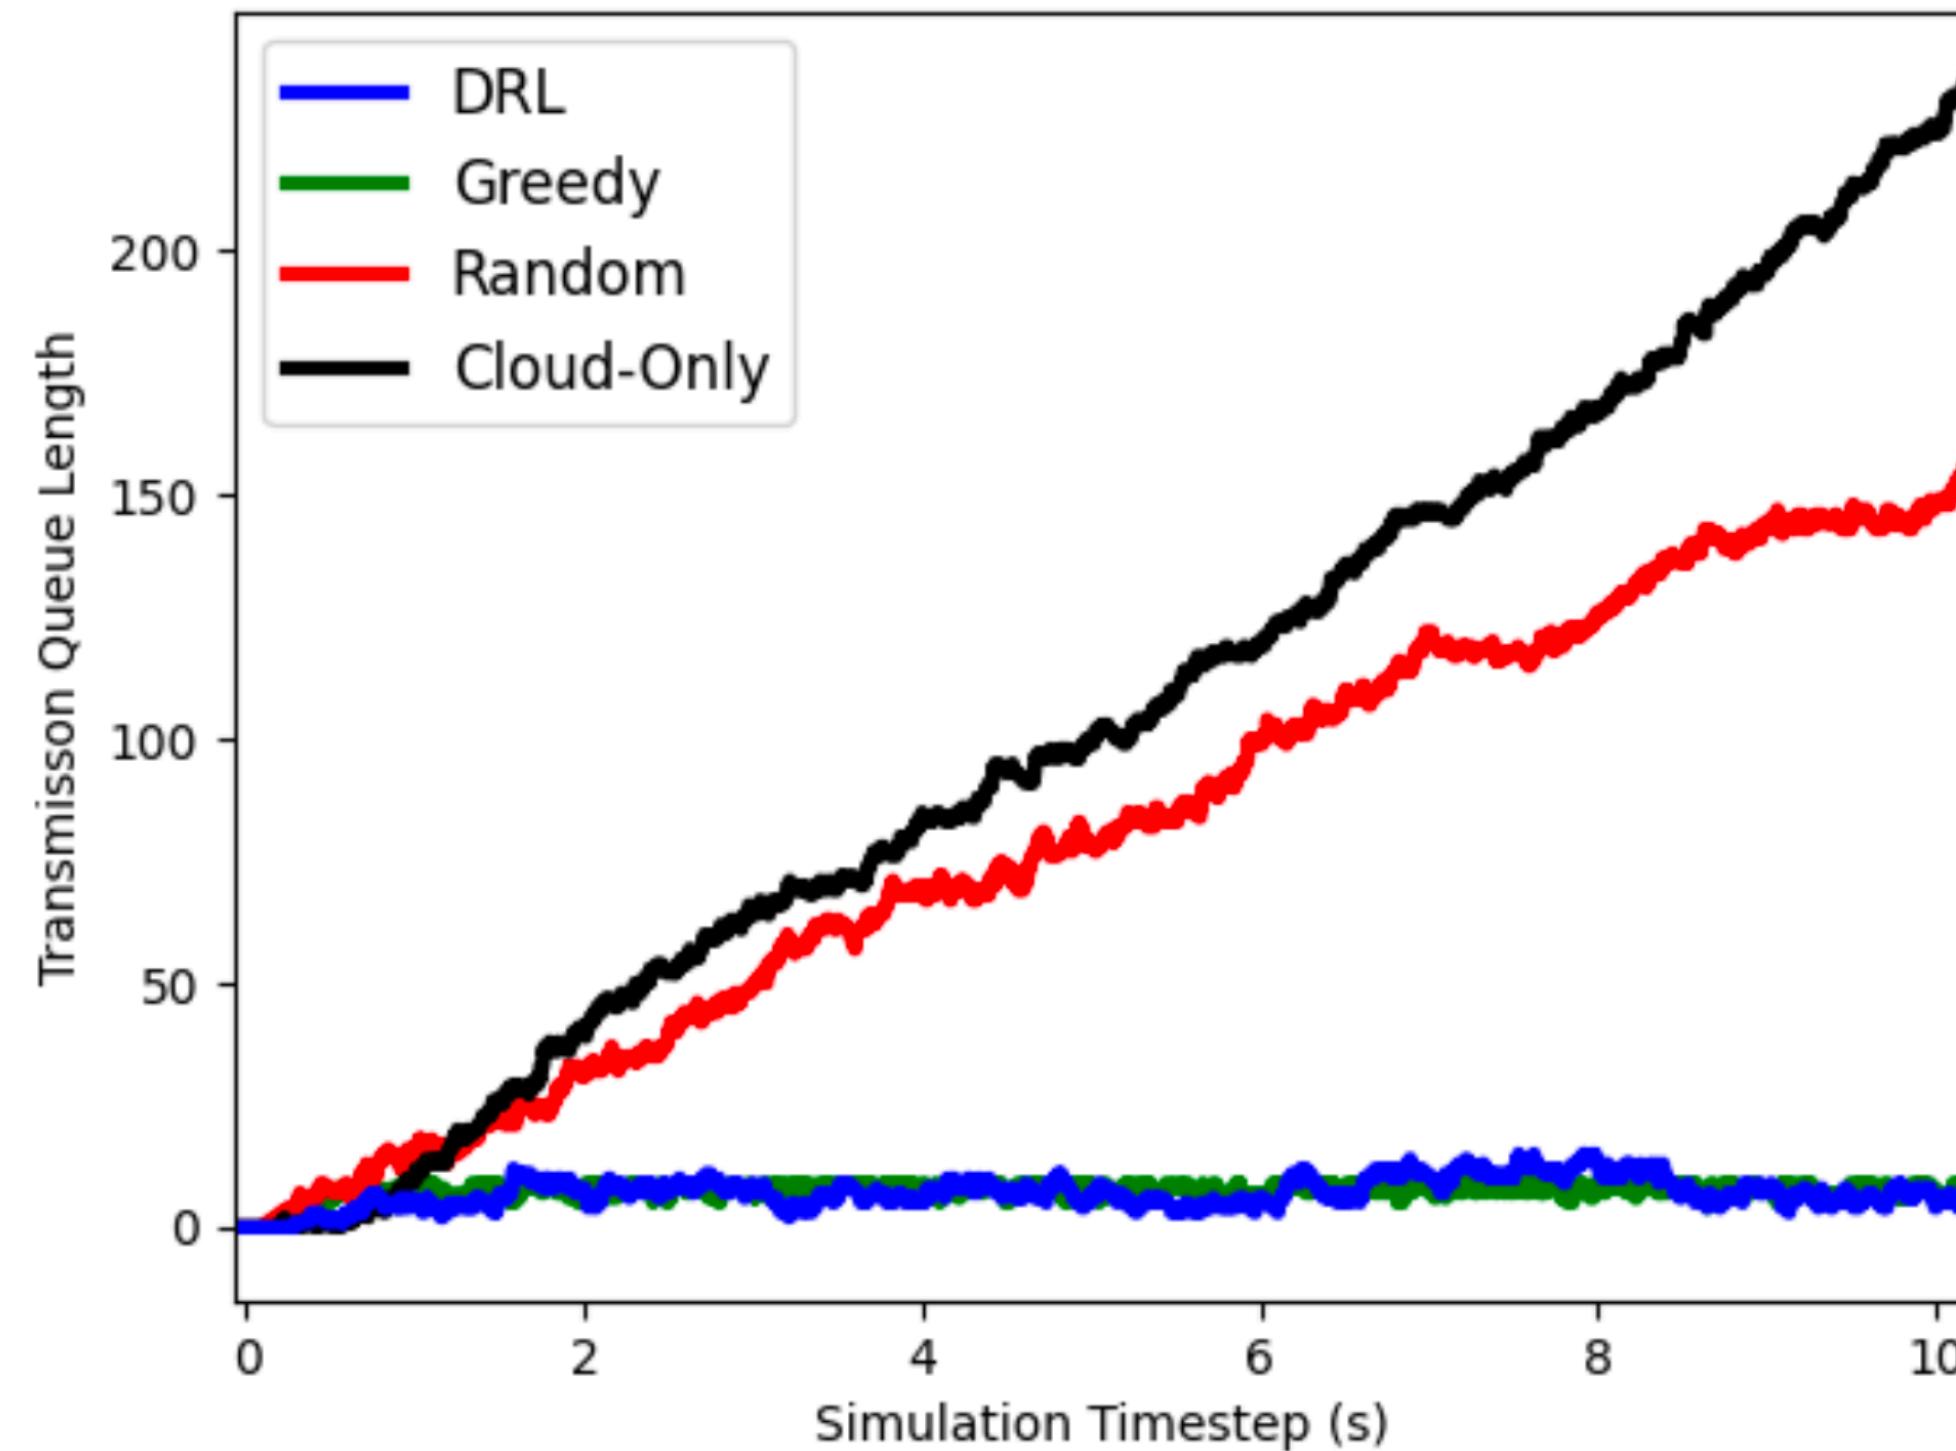
$$d_{\text{w-trans}} = \frac{CU_x}{TR_w}$$

TABLE I: Notation used for system model

Symbol	Meaning
l, w	Simulation area length & width (m)
n	Number of client vehicles
m	Number of service vehicles
CPU_j	Processing power of vehicle j (MIPS)
λ_i	Average task emission rate by client i
vel_i	Velocity of client i
CU_x	Number of million instructions of task x
SZ_x	Size of task x (Mb)
TR_w	Transmission rate of medium w (Mbps)
B	Wireless channel bandwidth (MHz)
N_0	Noise Power (dBm/Hz)
P_T	Transmitter's transmission Power (W)
G_T/G_R	Transmitting/Receiving directional gain (dBi)

Extra Slide #2

Transmission queue length



Traffic Intensity = $(aL)/R$

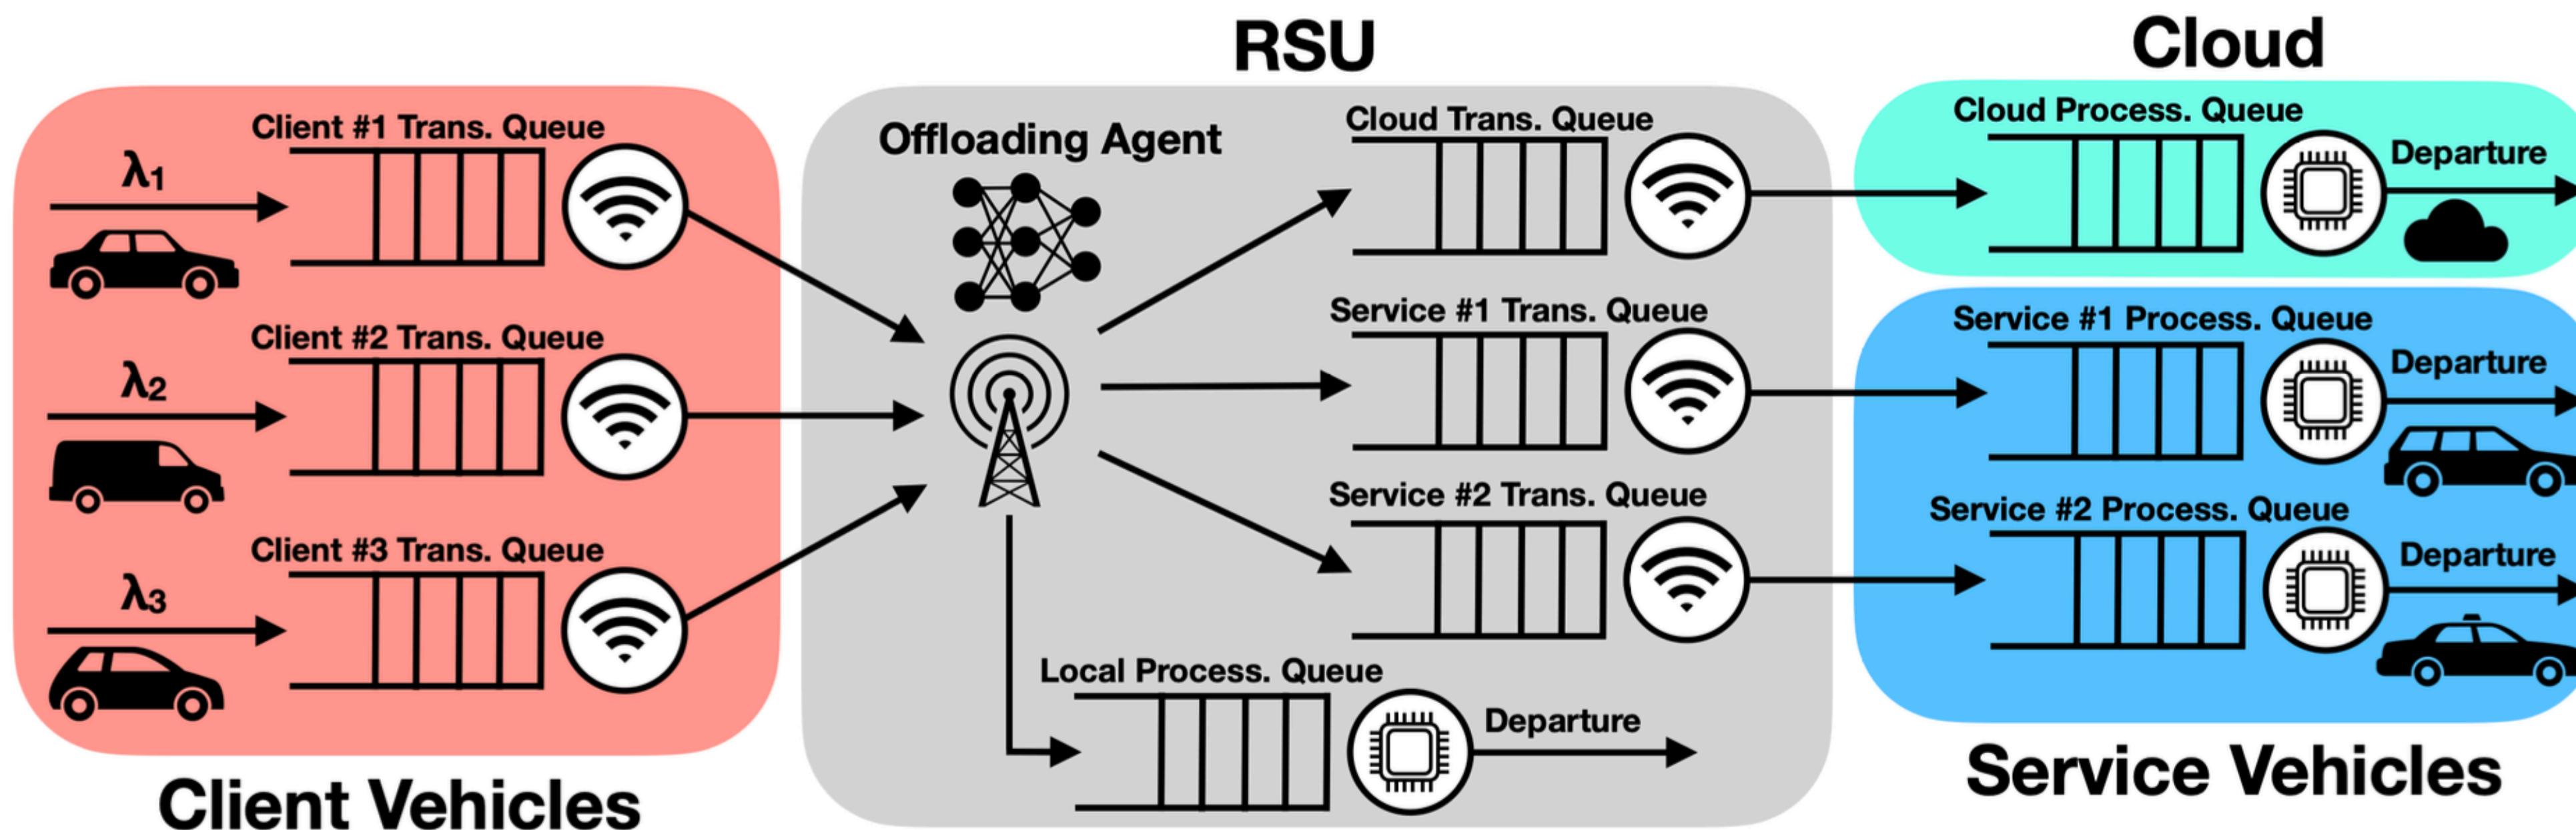
For cloud only: $(3 \times 20) \times 30 / 1000 = 1.8 > 1$

Fig. 6: Total length of all transmission queues throughout the first 10 seconds of simulation (Scenario 3).

Extra Slide #3

Delay & Communication model

- The total delay experienced by a task is the sum of the following: $d_{\text{total}} = d_{\text{client}} + d_{\text{rsu}} + d_{\text{service}}$



$$d_{w-\text{trans}} = \frac{CU_x}{TR_w}$$

$$TR = B \cdot \log_2 (1 + SNR)$$

$$SNR = \frac{P_R}{N_0 \cdot W}$$

$$P_R = P_T \cdot G_R \cdot G_T \cdot \left(\frac{\lambda}{4\pi d} \right)^2$$

Extra Slide #4

VFC Motivation

- Estimated 1.5 billion cars globally
- This number is projected to reach 2 billion by 2040
- Assume an average of 10 GFlops per vehicle
- $1.5 \text{ billion cars} * 10\text{GFlops} = 1500 \text{ Peta Flops}$
- The most powerful supercomputer system of 2024 has 1195 Peta Flops
-